

Modeling core financial domains with RESTful + CQRS at Nubank

Chapter 1: Domain

Avoiding Objects, using Values¹

¹ https://www.youtube.com/watch?v=ROor6_NGIWU

Introduction: What is a loan?

A loan is the lending of money (...)

The borrower incurs a debt, and is usually liable to pay interest on that debt until it is repaid, and also to repay the principal amount borrowed.⁰

- What we do at Nubank
- We better get this right :)

⁰ <https://en.wikipedia.org/wiki/Loan>

Idea: Use Business Language

A language structured around the **domain model** and **used by all team members** to connect all the activities of the team with the software.²

² https://en.wikipedia.org/wiki/Domain-driven_design

Domain-Driven DESIGN

Tackling Complexity in the Heart of Software



Eric Evans
Foreword by Martin Fowler

What is a loan?

- Principal (\$)
- Interest Rate (%)
- Late Fee (%)
- Maturity (Date)
- Payment schedule (Bullet, Straight-line, Constant Amortization, ...)
- Purpose (Credit Card, Personal Loan, Mortgage, ...)

What is a loan?

```
#:loan{:id          UUID
       :type        #{:loan.type/personal, :loan.type/credit-card, :loan.type/mortgage, #_...}
       :amortization-schedule #{:loan.amortization-schedule/bullet, :loan.amortization-schedule/price, #_...}
       :installments Integer
       :contract-date LocalDate
       :interest-rate BigDecimal
       :late-interest-rate BigDecimal
       :late-fee      BigDecimal
; ...
}
```

- Not a class, not an object
- No behaviour implied, no state
- Just data

What can happen to a loan?

- Payment
- Late payment
- Anticipation
- Renegotiation
- Cancellation
- "Write-off" (credit default, fraud, ...)

Idea: Event Sourcing

Capture all **changes** to an application state as a **sequence of events**.⁵

⁵ <https://martinfowler.com/eaaDev/EventSourcing.html>

What can happen **to** a loan?

```
#loan-event{  
    :id          UUID  
    :type        #{:loan-event.type/payment,  
                  :loan-event.type/late,  
                  :loan-event.type/anticipation, #_...}  
    :reference-date LocalDate  
    :sort-order    Integer  
    :disbursement  BigDecimal  
    :discount      BigDecimal  
    :interest      BigDecimal  
    :late-fee       BigDecimal  
    :late-interest  BigDecimal  
    ; ...  
}
```

- Remember **what** happened, not the interpretation
- Explicit date to be relevant, ordered
- Can be stored, queued, logged...

What can happen **to** a loan?

(handle-event initial-view issue-event) => view'

(handle-event view' payment-event) => view''

(reduce handle-event initial-view [event ...]) => view'''

- Rely on **pure functions** and **higher-order functions**
- Focus on computation rather than updating objects
- Easy to **observe, test**

What can happen **to** a loan?

```
(defmethod snapshot-after :loan-event.type/due-payment
  [snapshot-before new-event]
  (let [#_...compute a lot stuff...]
    (-> snapshot-before
        (assoc :debt-snapshot/pv-without-late pv-at-due-date)
        (assoc :debt-snapshot/next-due-date (some-> next-charge :due-date))
        (assoc :debt-snapshot/accruals accruals-after-event)
        (update :debt-snapshot/past-events conj event)
        (update :debt-snapshot/future-charges (comp vec rest))
        (update :debt-snapshot/paid-charges conj paid-charge))))
```

```
1 #:debt-snapshot{:past-events [#:loan-event{:reference-date #nu/date "2014-04-01"
2 :loan #:loan{:id #uuid "5d586283-14bf-4ed2-a517-456bf0f3194c"}
3 :type :loan-event.type/issuance
4 :disbursement 1000M
5 :fixed-iof 3.85M
6 :daily-iof 9.64M
7 :financed-iof 13.49M
8 :id #uuid "5d586283-248d-4feb-8517-f22ac5ad65cc"}
9 #:loan-event{:id #uuid "5d5861a2-126d-447a-99c6-3d230fb634aa"
10 :type :loan-event.type/late
11 :reference-date #nu/date "2014-05-01"
12 :loan #:loan{:customer-id #uuid "5d58619b-6662-4416-b773-14a795e"
13 :interest-rate 0.11M
14 :yearly-cet 0.15M
15 :installments 6
16 :late-fee 0.02M
17 :type :loan.type/personal
18 :due-day 1
19 :contract-date #nu/date "2014-04-01"
20 :status :loan.status/pending
21 :id #uuid "5d58619b-1b73-4bdb-a44e-1651e5e"
22 :initial-days-late 0
23 :late-interest-rate 0.12M
24 :issuer :loan.issuer/nubank-financeira
25 :holder :loan.holder/nubank-financeira
26 :first-due-date #nu/date "2014-05-01"
27 :principal 1000M
28 :grace-days 30
29 :amortization-schedule :loan.amortization-schedule/price}
30 :charge-index 0}
31 #:loan-event{:late-fee -4.7913056686M
32 :complementary-iof -0.32558288M
33 :flow-id "banana:late-payment"
34 :type :loan-event.type/late-payment
35 :loan #:loan{:customer-id #uuid "5d58619b-6662-4416-b773-14a7
36 :interest-rate 0.11M
37 :yearly-cet 0.15M
38 :installments 6
39 :late-fee 0.02M
40 :type :loan.type/personal
41 :due-day 1
42 :contract-date #nu/date "2014-04-01"
```

Idea: Persistent Data Structures⁴

⁴ <https://www.infoq.com/presentations/Value-Identity-State-Rich-Hickey/>

Persistent Data Structures and Managed References

LIKE

9

BOOKMARKS



View Presentation



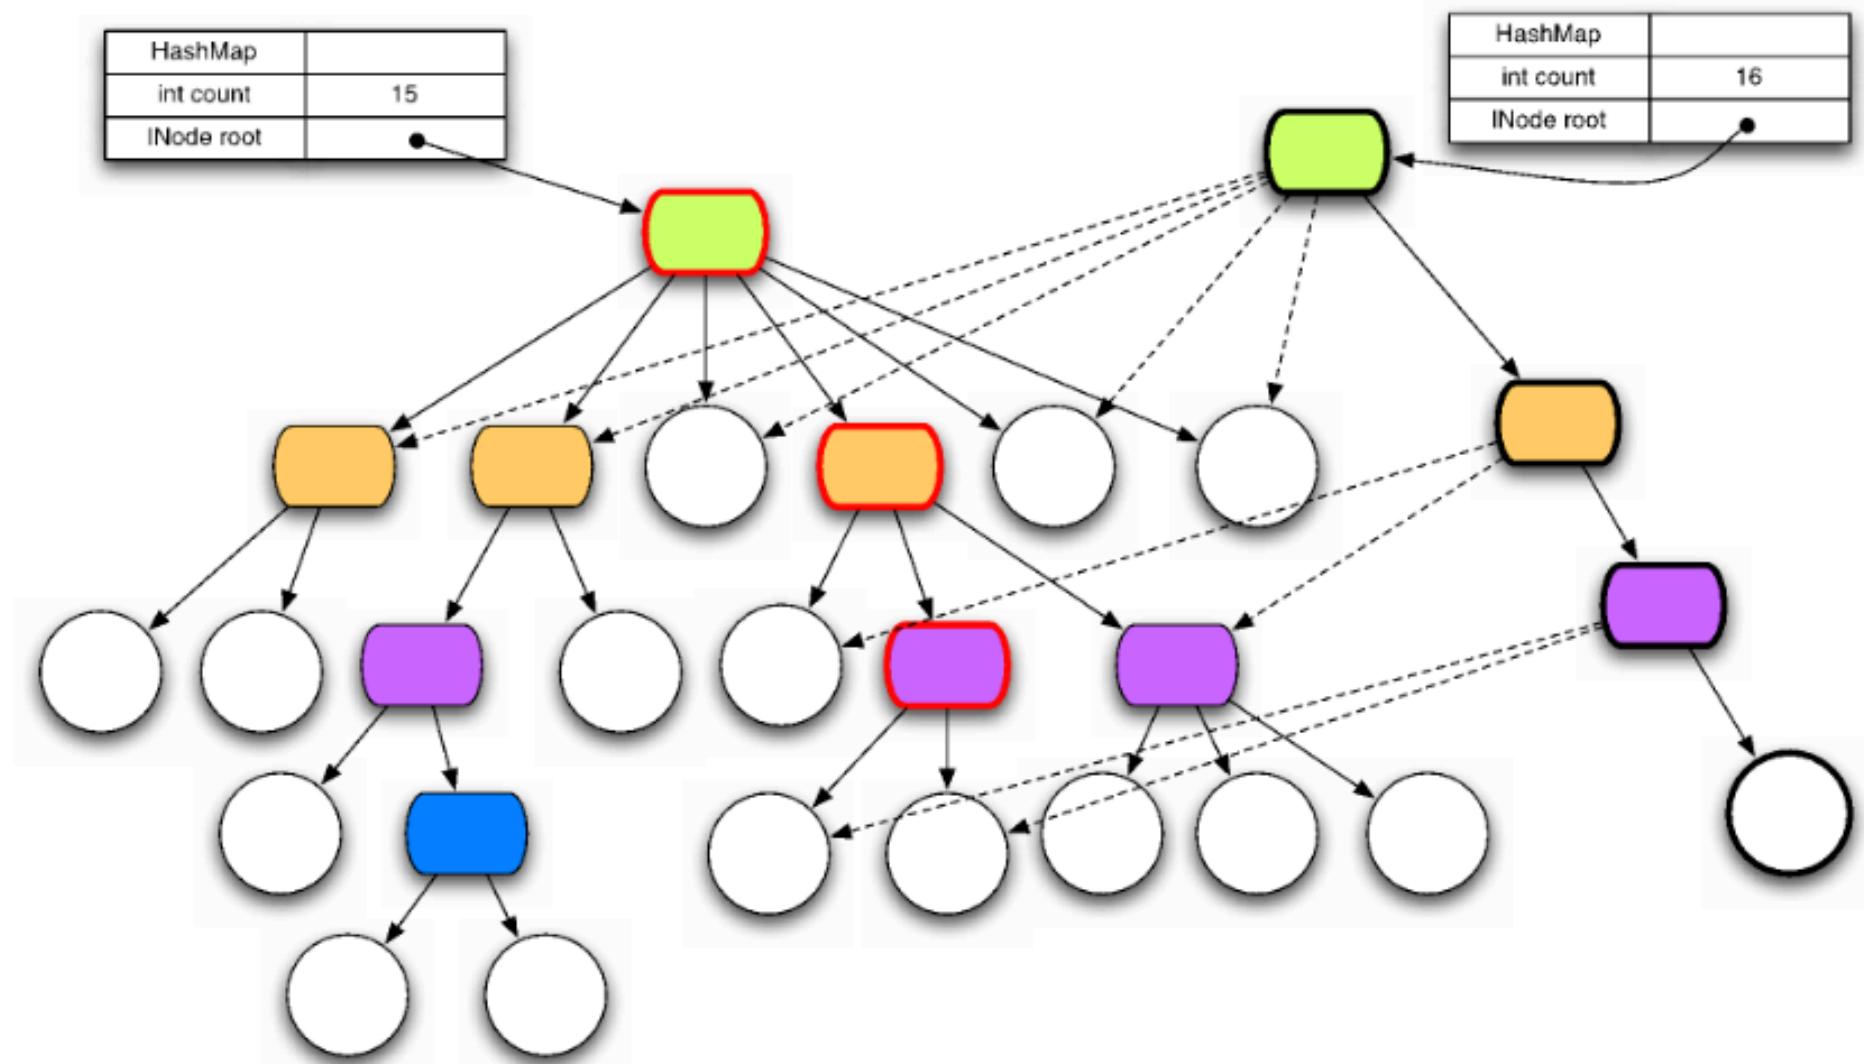
Speed: 1X 1.5X 2X



Summary

Rich Hickey's presentation is organized around a number of programming concepts: identity, state and values. He explains how to represent composite objects as values and how to deal with change and state, as it is implemented in Clojure.

Path Copying



Bonus: easy parallelism

Parallelize issuance simulations to decrease endpoint latency
#432

Edit

Merged hcarvalhoalves merged 1 commit into master from parallelize-simulations on Jan 11

Conversation 1 Commits 1 Checks 0 Files changed 1 +1 -1

Changes from all commits ▾ File filter... ▾ Jump to... ▾ 0 / 1 files viewed Review changes ▾

src/corleone/controllers/lending.clj

@@ -142,6 +142,6 @@	Viewed	...
142 142		
143 143 (s/defn issuance-simulations :- (Either models.error/CantSimulate [models.debt/DebtSnapshot])		
144 144 [commands :- [models.command/Issuance], db :- Db]		
145 - (->> (map #(simulate % db) commands)		
145 + (->> (pmap #(simulate % db) commands)		
146 146 cats/sequence		
147 147 (cats/fmap (partial map :current-debt))))		

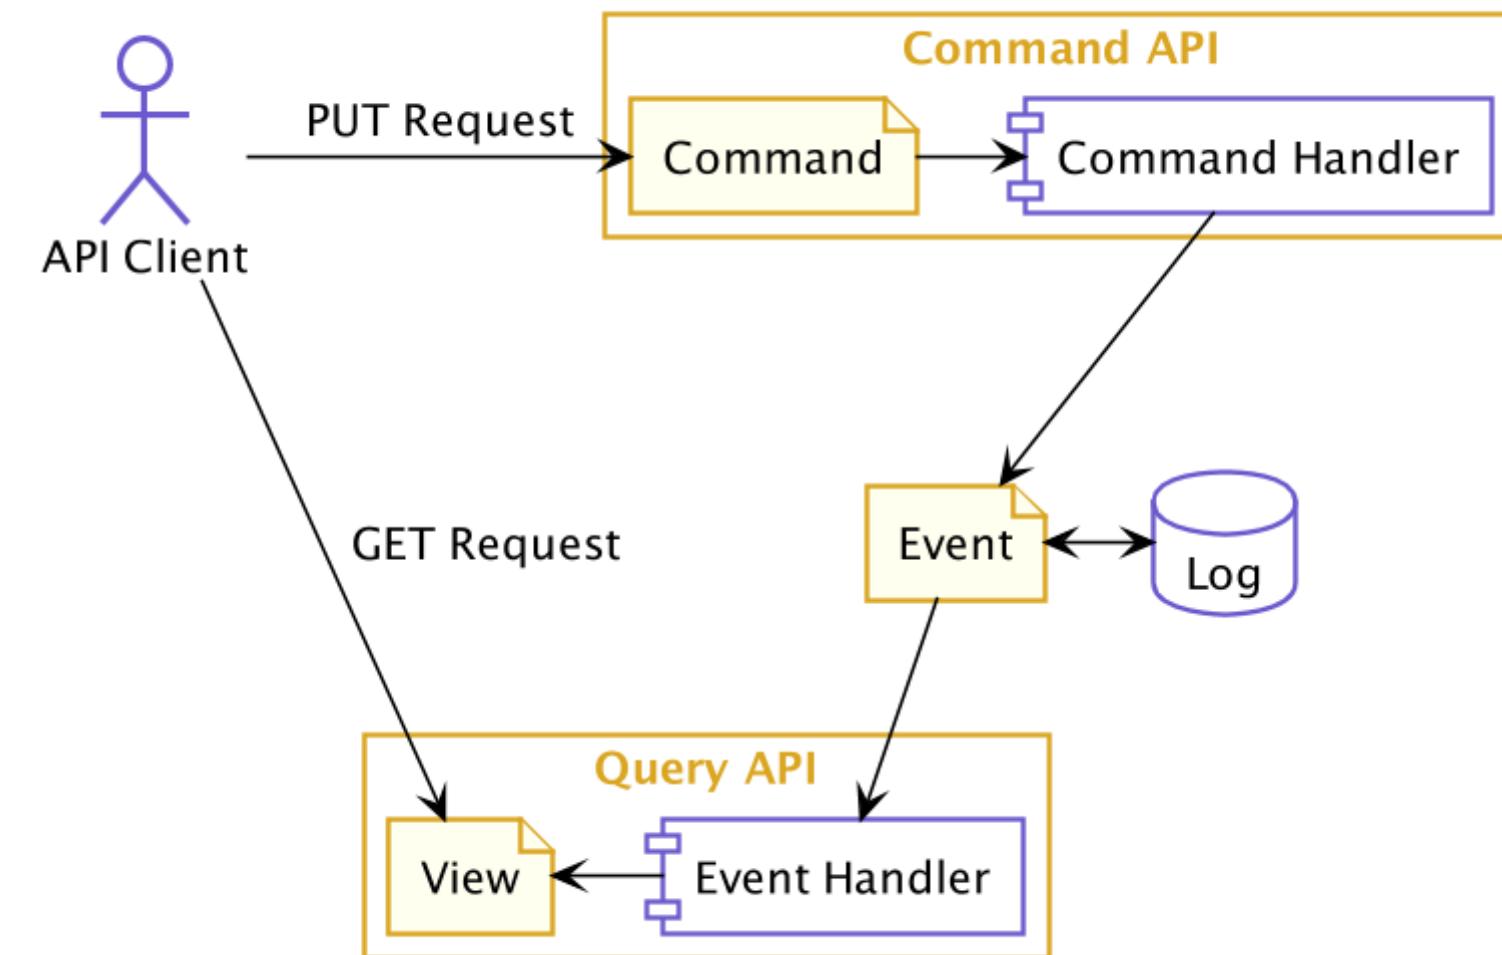
What causes updates to a loan?

- Customer borrows a loan
- Customer requests a payment
- Billing system requests an automatic payment
- Billing system detects the loan is late
- Fraud system detects a fraud

Idea: Command / query Responsibility Segregation

At its heart is the notion that you can use a **different model** to **update** information than the model you use to **read** information.⁶

⁶ <https://martinfowler.com/bliki/CQRS.html>



What causes updates to a loan?

```
#:command{:id
           :type      UUID
           #:command.type/issuance-request,
           #:command.type/payment-request,
           #:command.type/cancellation-request,
           #_...}
:requested-at    LocalDateTime
:processed-at    LocalDateTime
:reference-date LocalDate
:processing-status #:command.status/accepted, #:command.status/ok, #:command.status/failed, #_...}
```

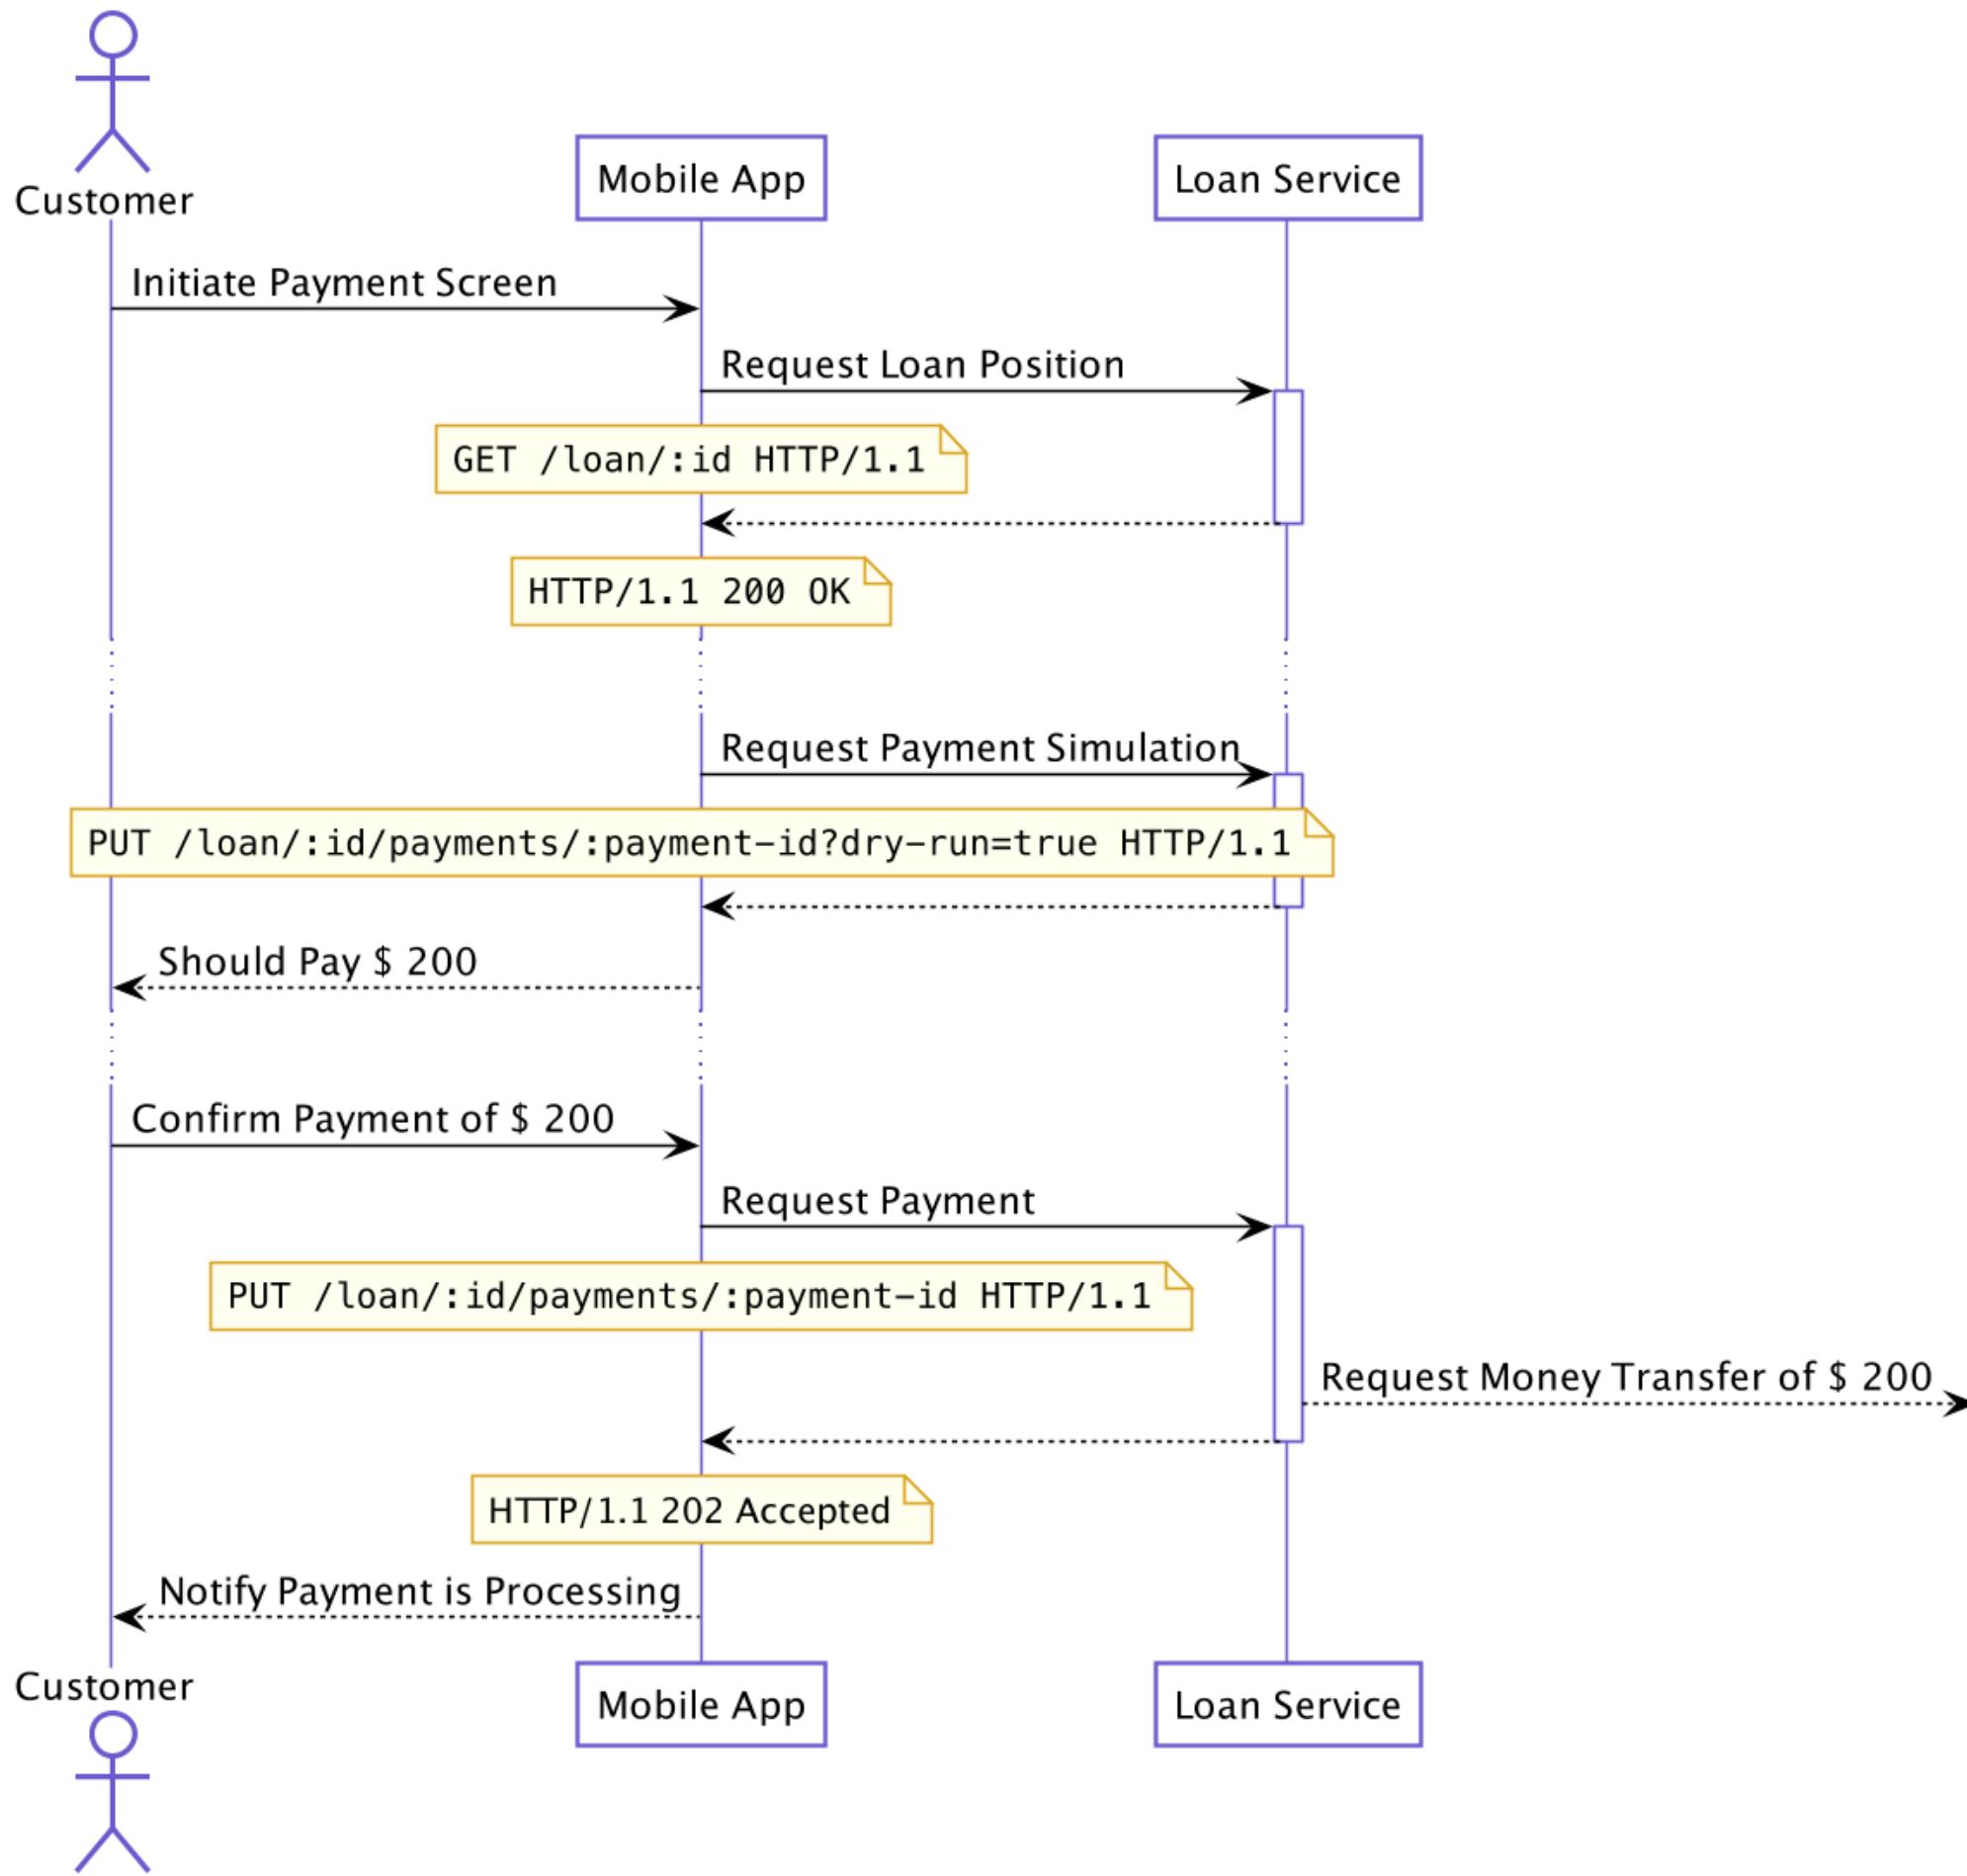
- Remember the **request**, assign a **unique id**
- Track status for **asynchronous** processing
- Can be stored, queued, retried, logged...

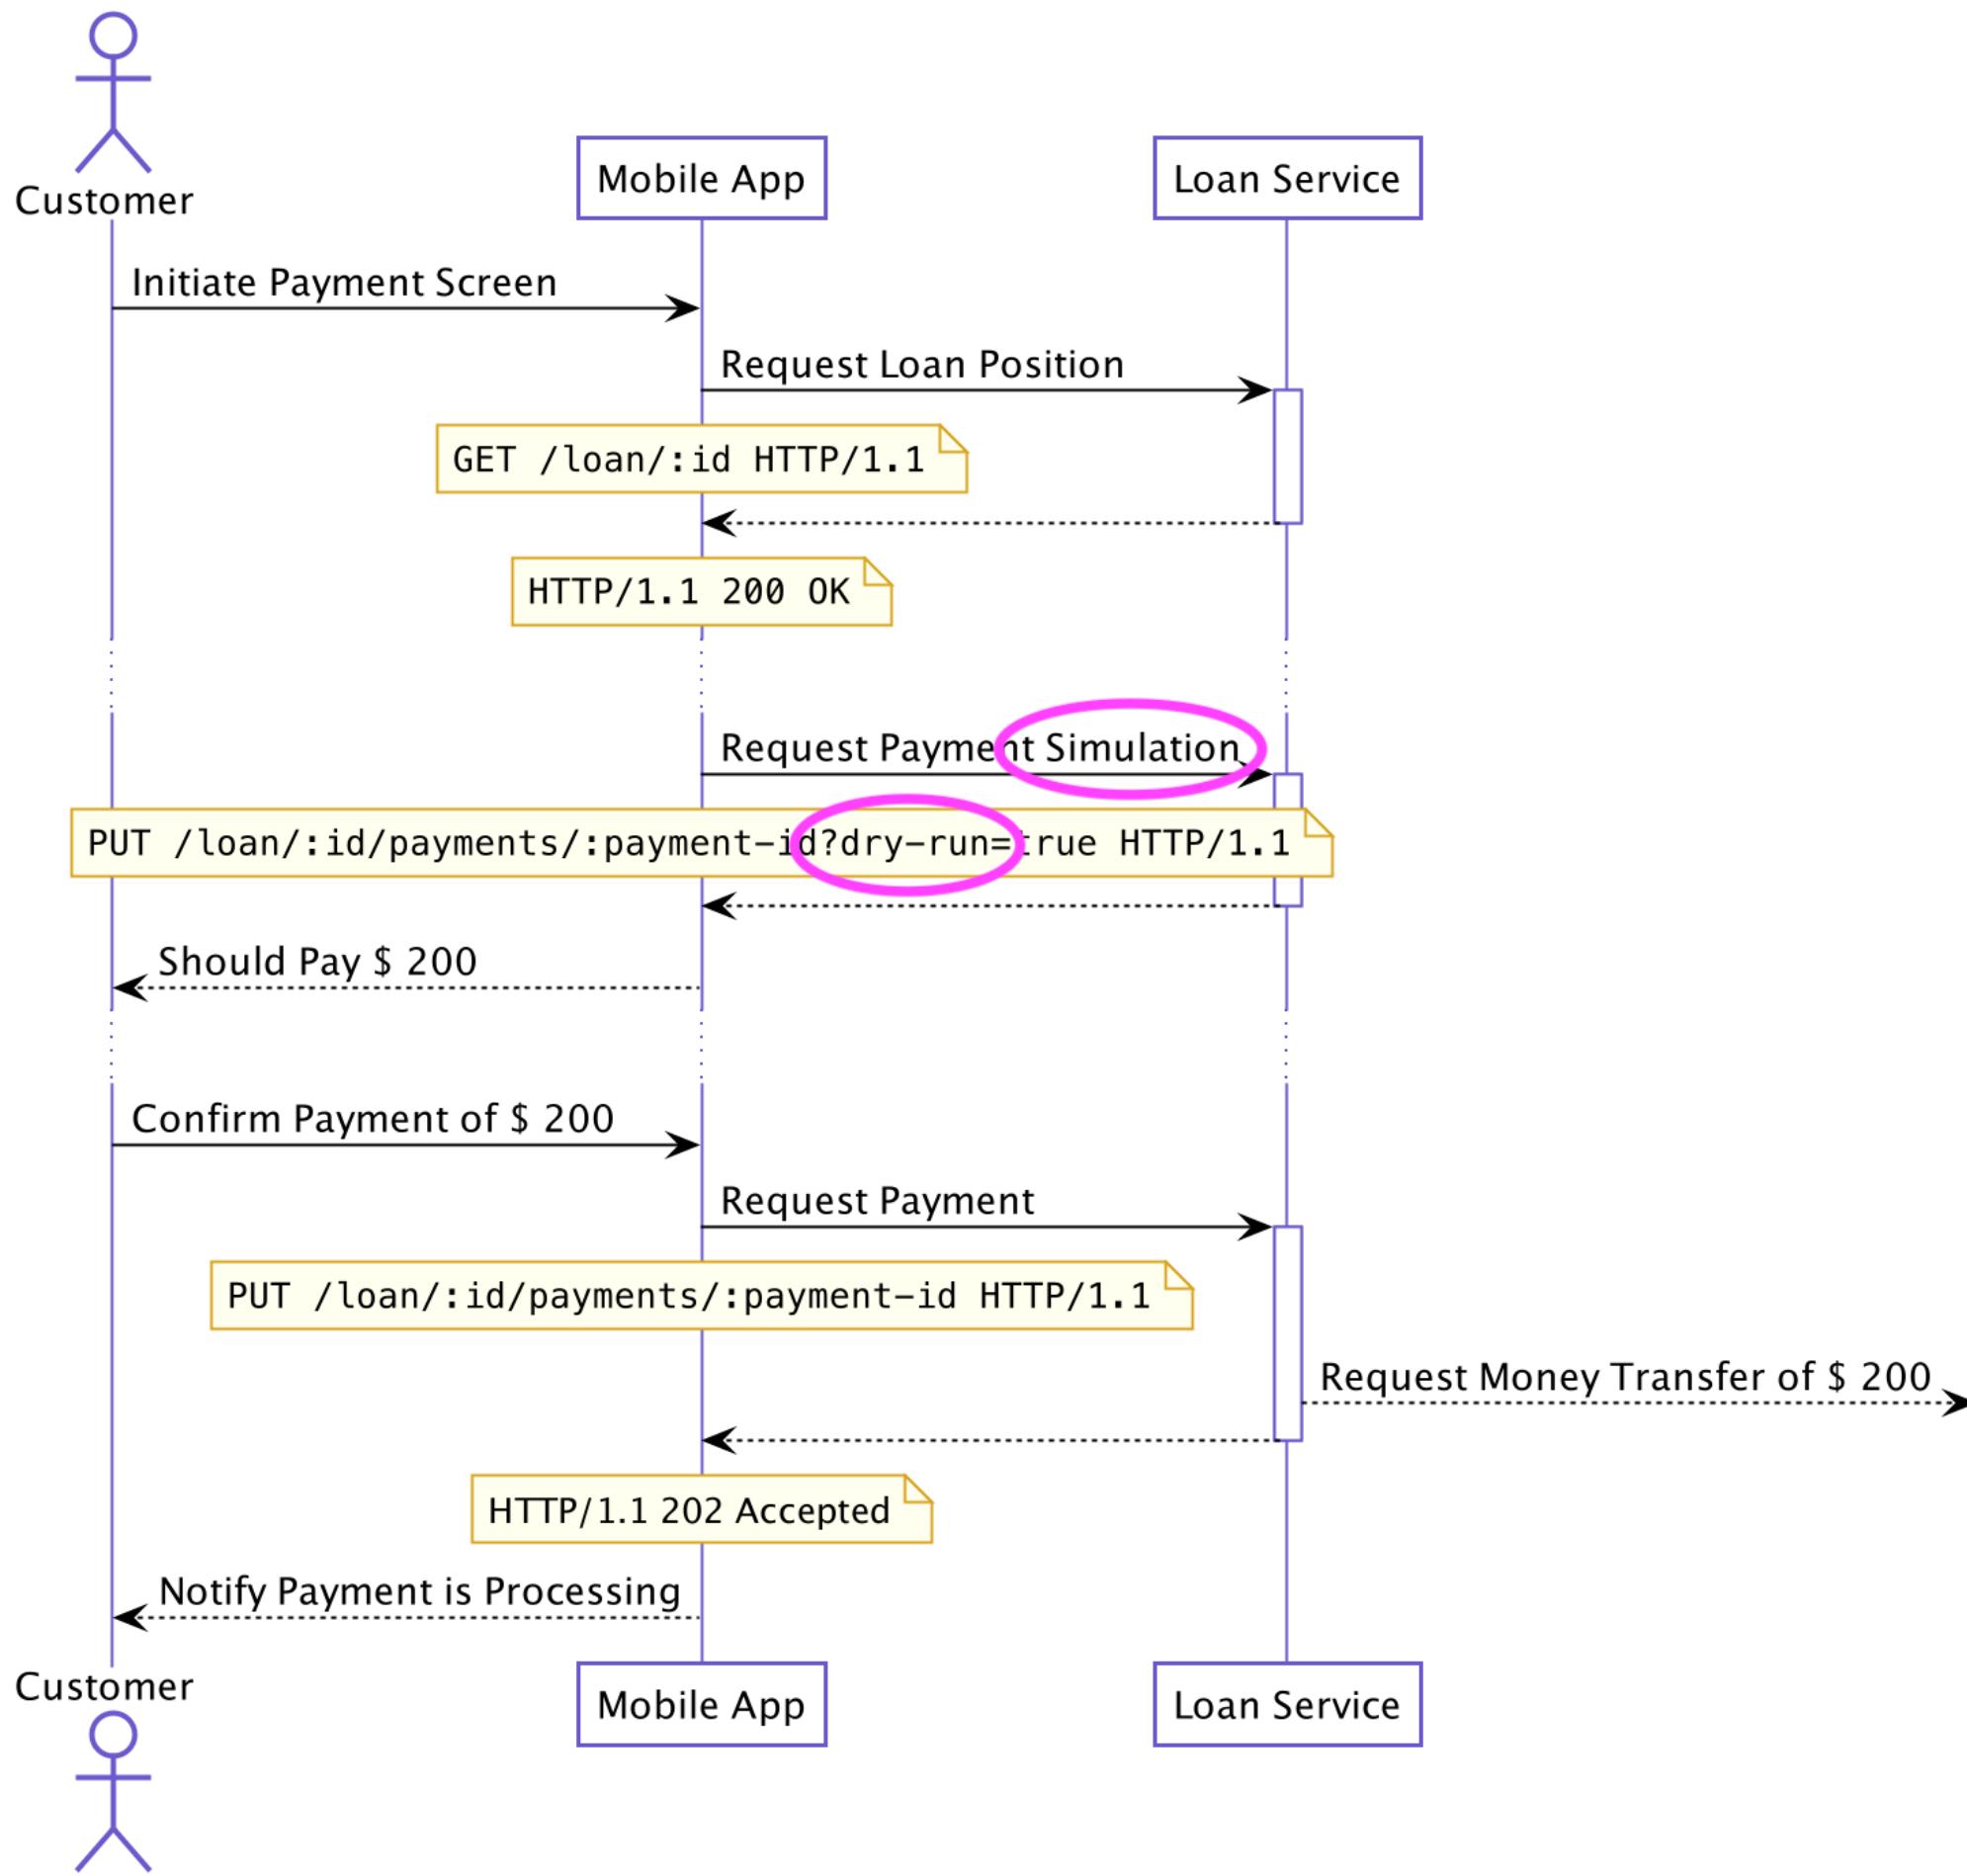
chapter 2: Flow

Flows vs. Places¹

¹ https://www.youtube.com/watch?v=ROor6_NGIWU

Practical example: Designing a Payment API





Idea: Datomic speculative transactions

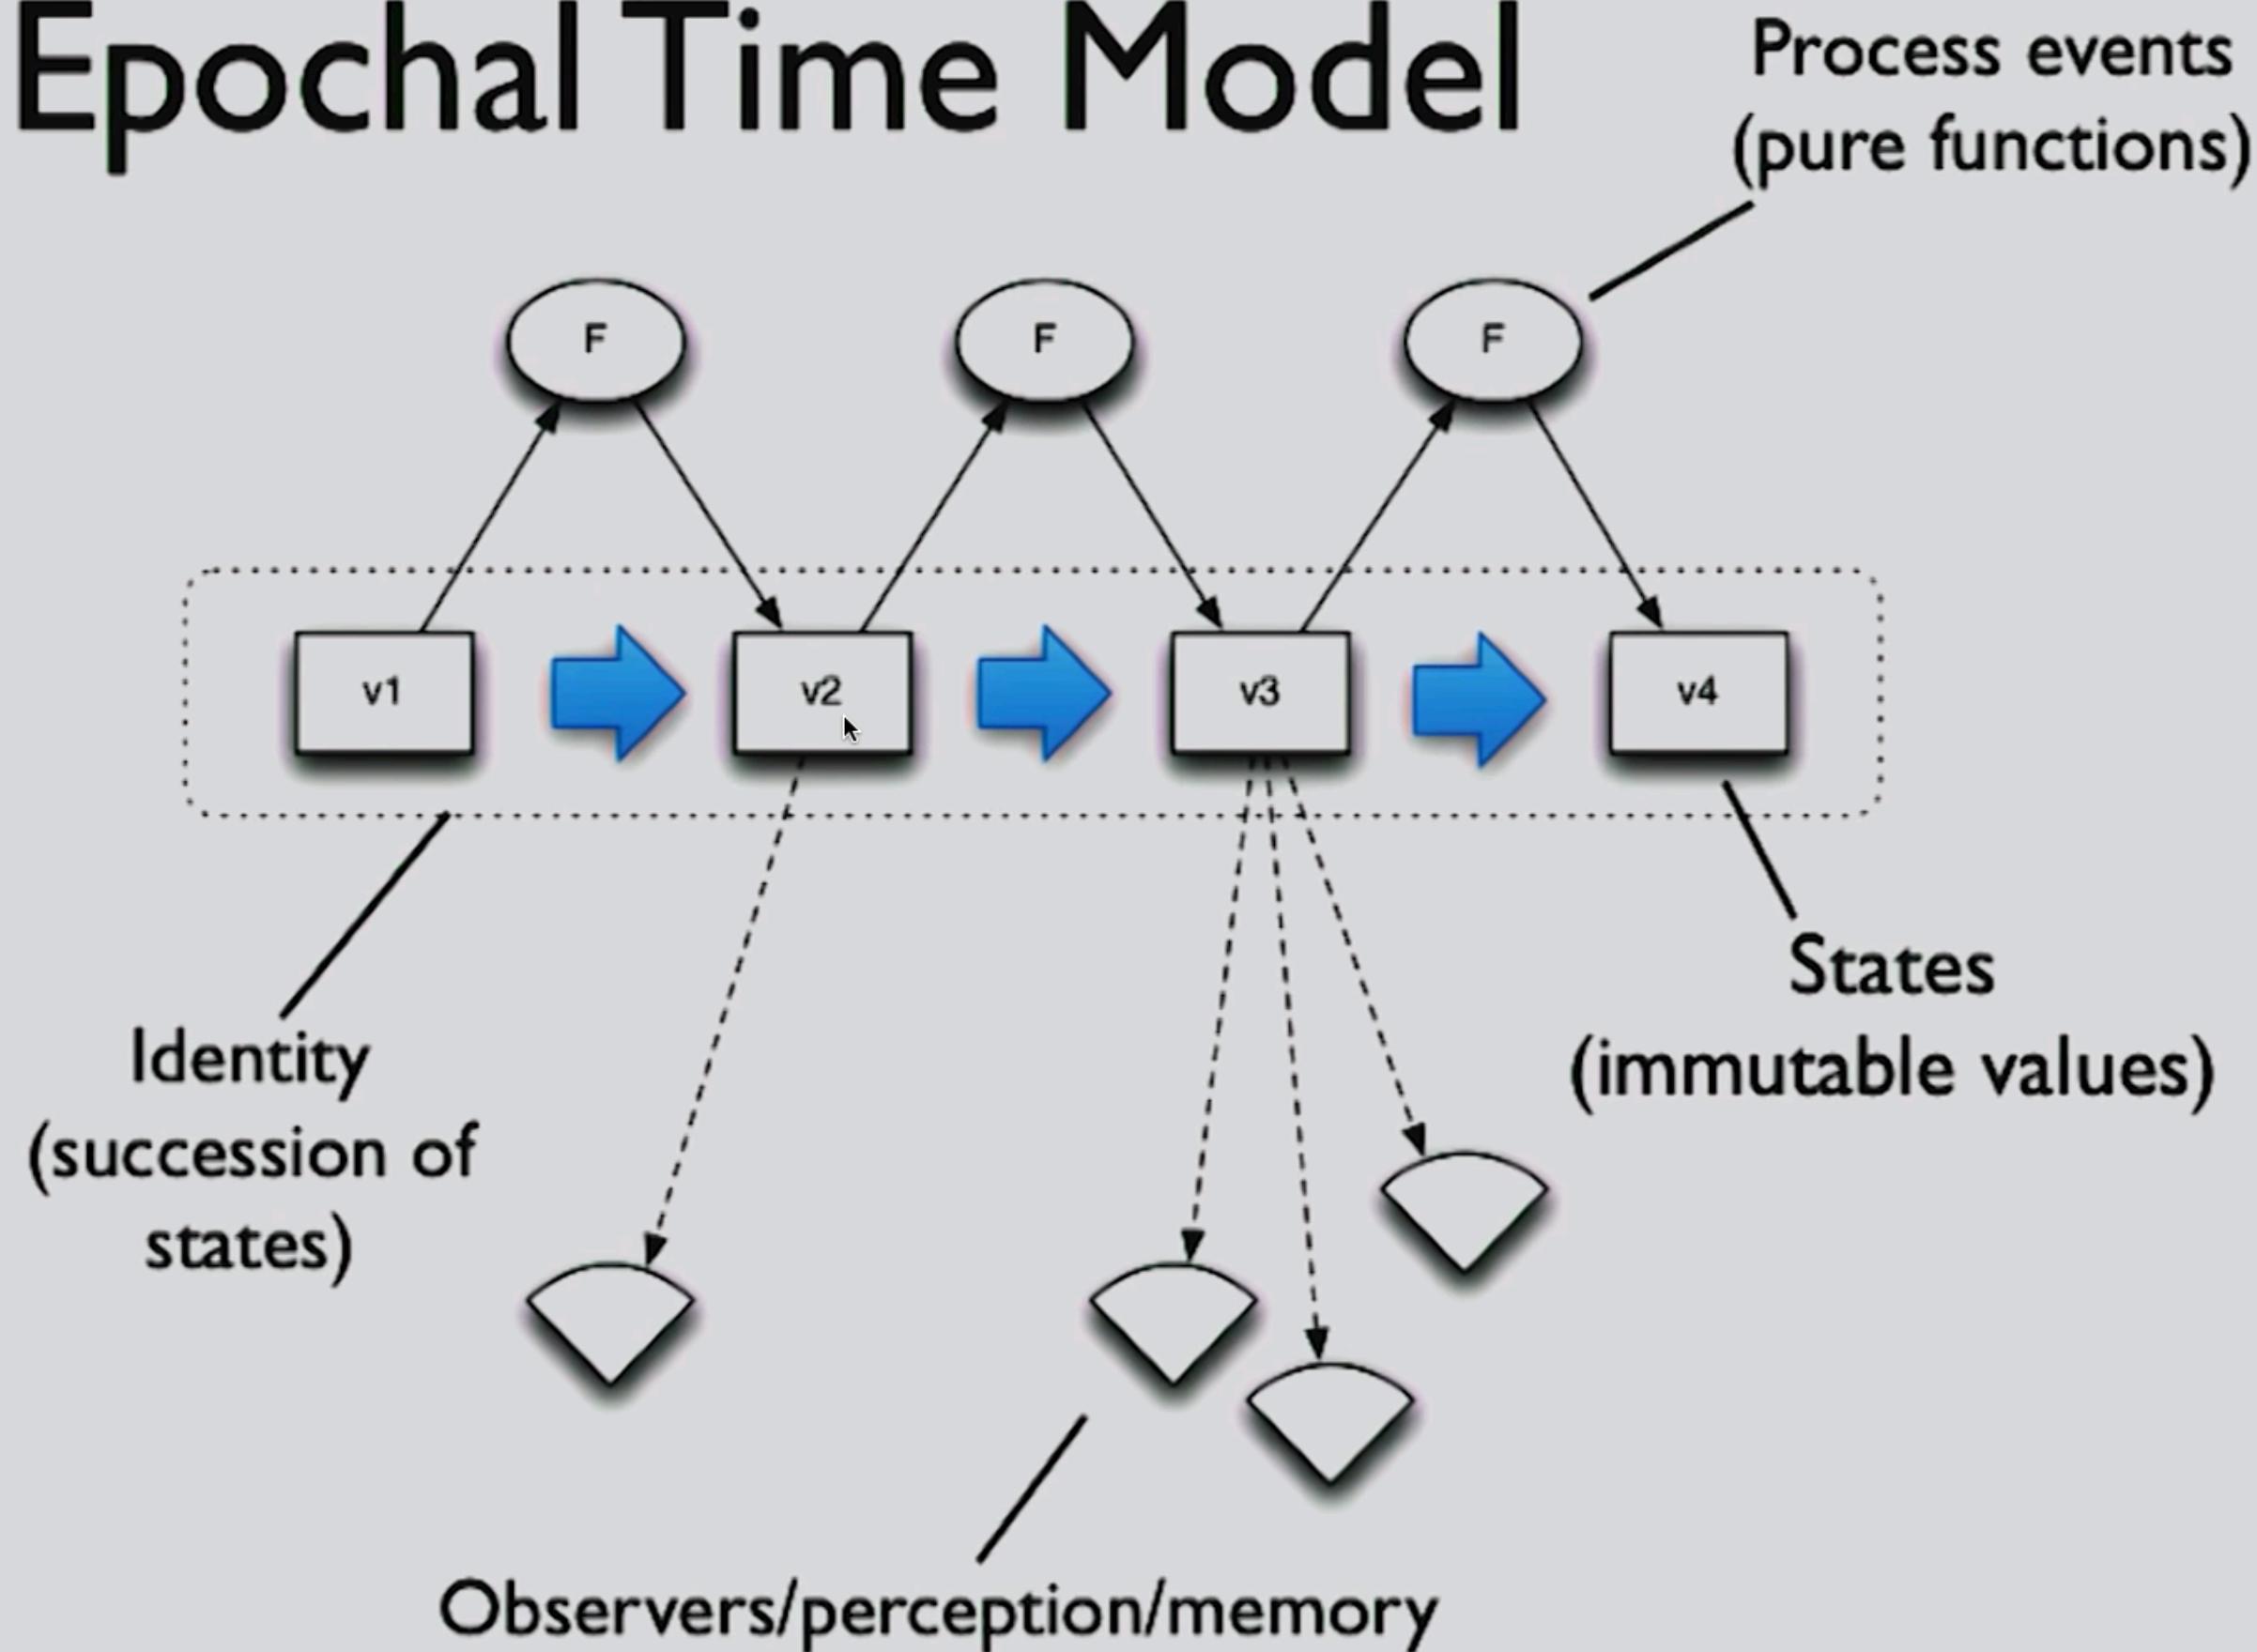


2012



Sheraton
Raleigh

Epochal Time Model



Quick Detour: Datomic's data model

entity

attribute

value

tx

Quick Detour: Datomic's data model

```
[{:name "Jane" :phone "12341234"}]
```

entity	attribute	value	tx
customer	name	Jane	1
customer	phone	12341234	1

Quick Detour: Datomic's data model

```
[{:name "Jane" :phone "912341234"}]
```

entity	attribute	value	tx
customer	name	Jane	1
customer	phone	12341234	1
customer	phone	912341234	2

Quick Detour: Datomic's data model

```
[{:name "Jane" :phone "912341234"}  
{:name "Joe"}]
```

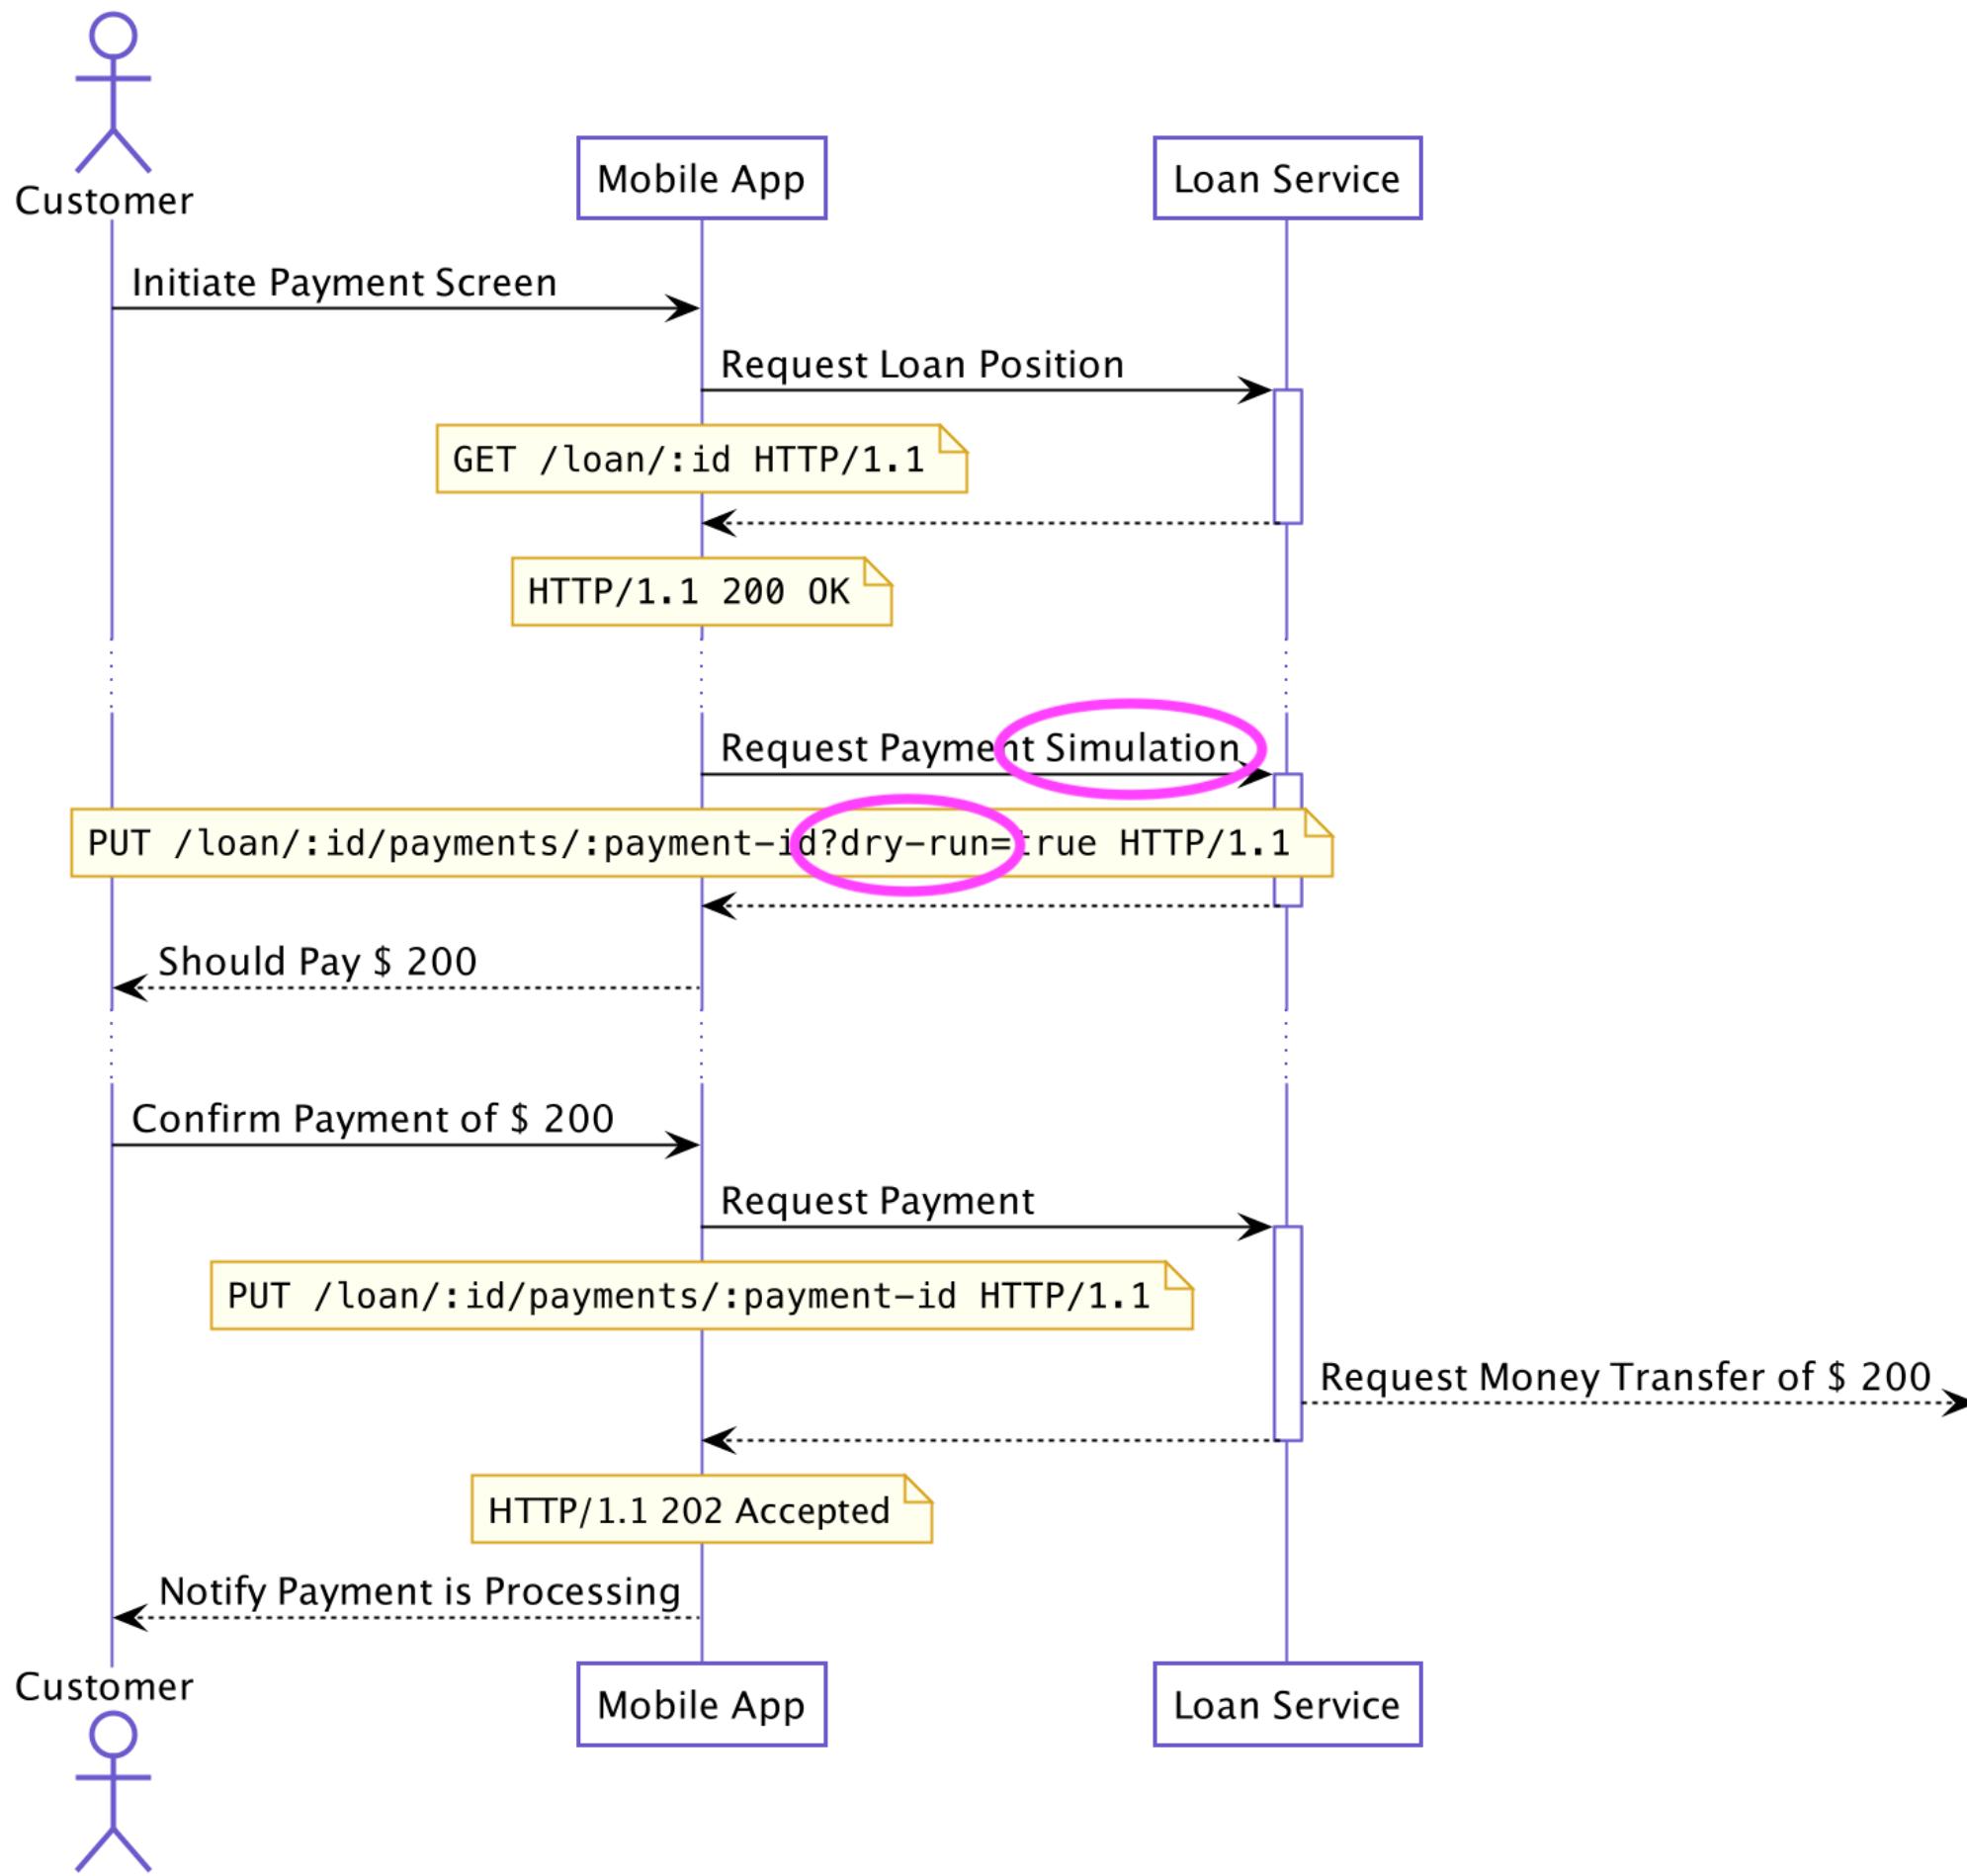
entity	attribute	value	tx
customer	name	Jane	1
customer	phone	12341234	1
customer	phone	912341234	2
customer2	name	Joe	3

Datomic speculative transactions

```
(let [events (command->events command)
      tx-data (prepare-transaction events)]
  (-> (if dry-run?
         (datomic.api/with db tx-data) ;; applies in-memory
         (datomic.api/transact datomic tx-data)) ;; persists to storage
       (get :db-after) ;; lazy representation of entire db
       (proceed-as-usual)))
```

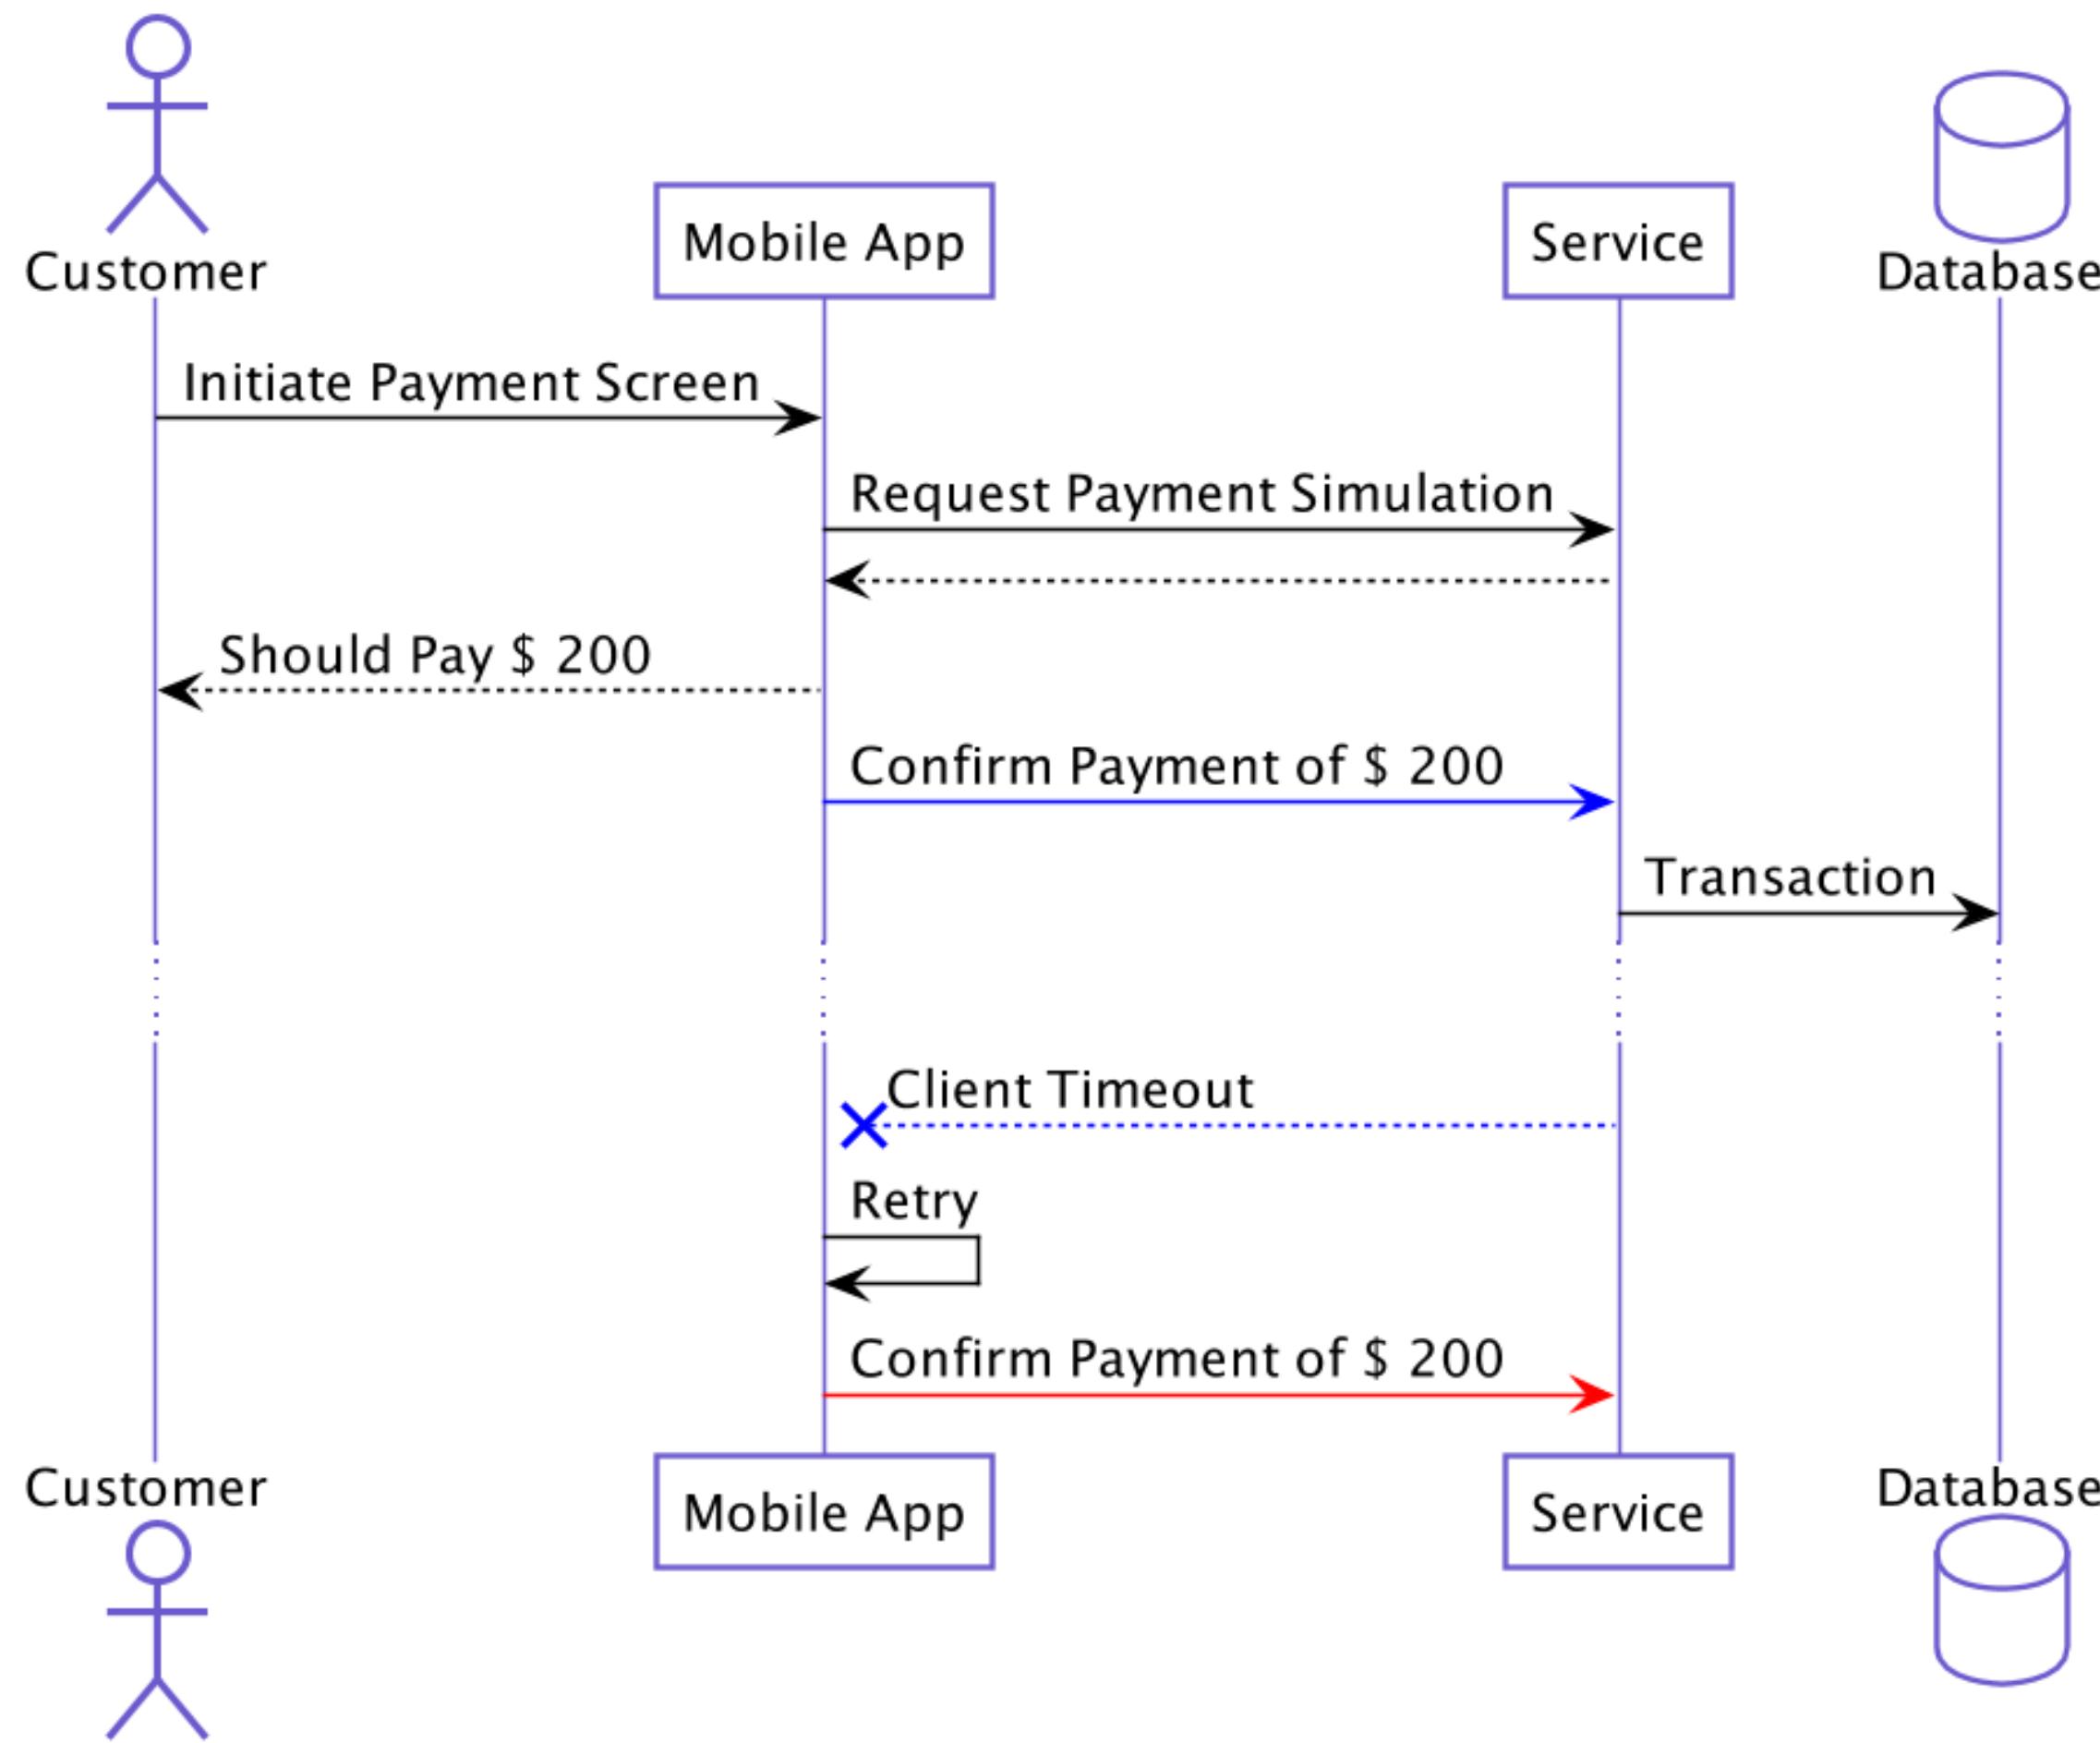
Datomic speculative transactions

- Database is handled as a (lazy) value
 - Pure functions
- Same code path for simulation or transaction
 - Reduced bugs by avoiding corner cases



Problem:

Idempotent updates



Idea: Datomic's historical database

as-of

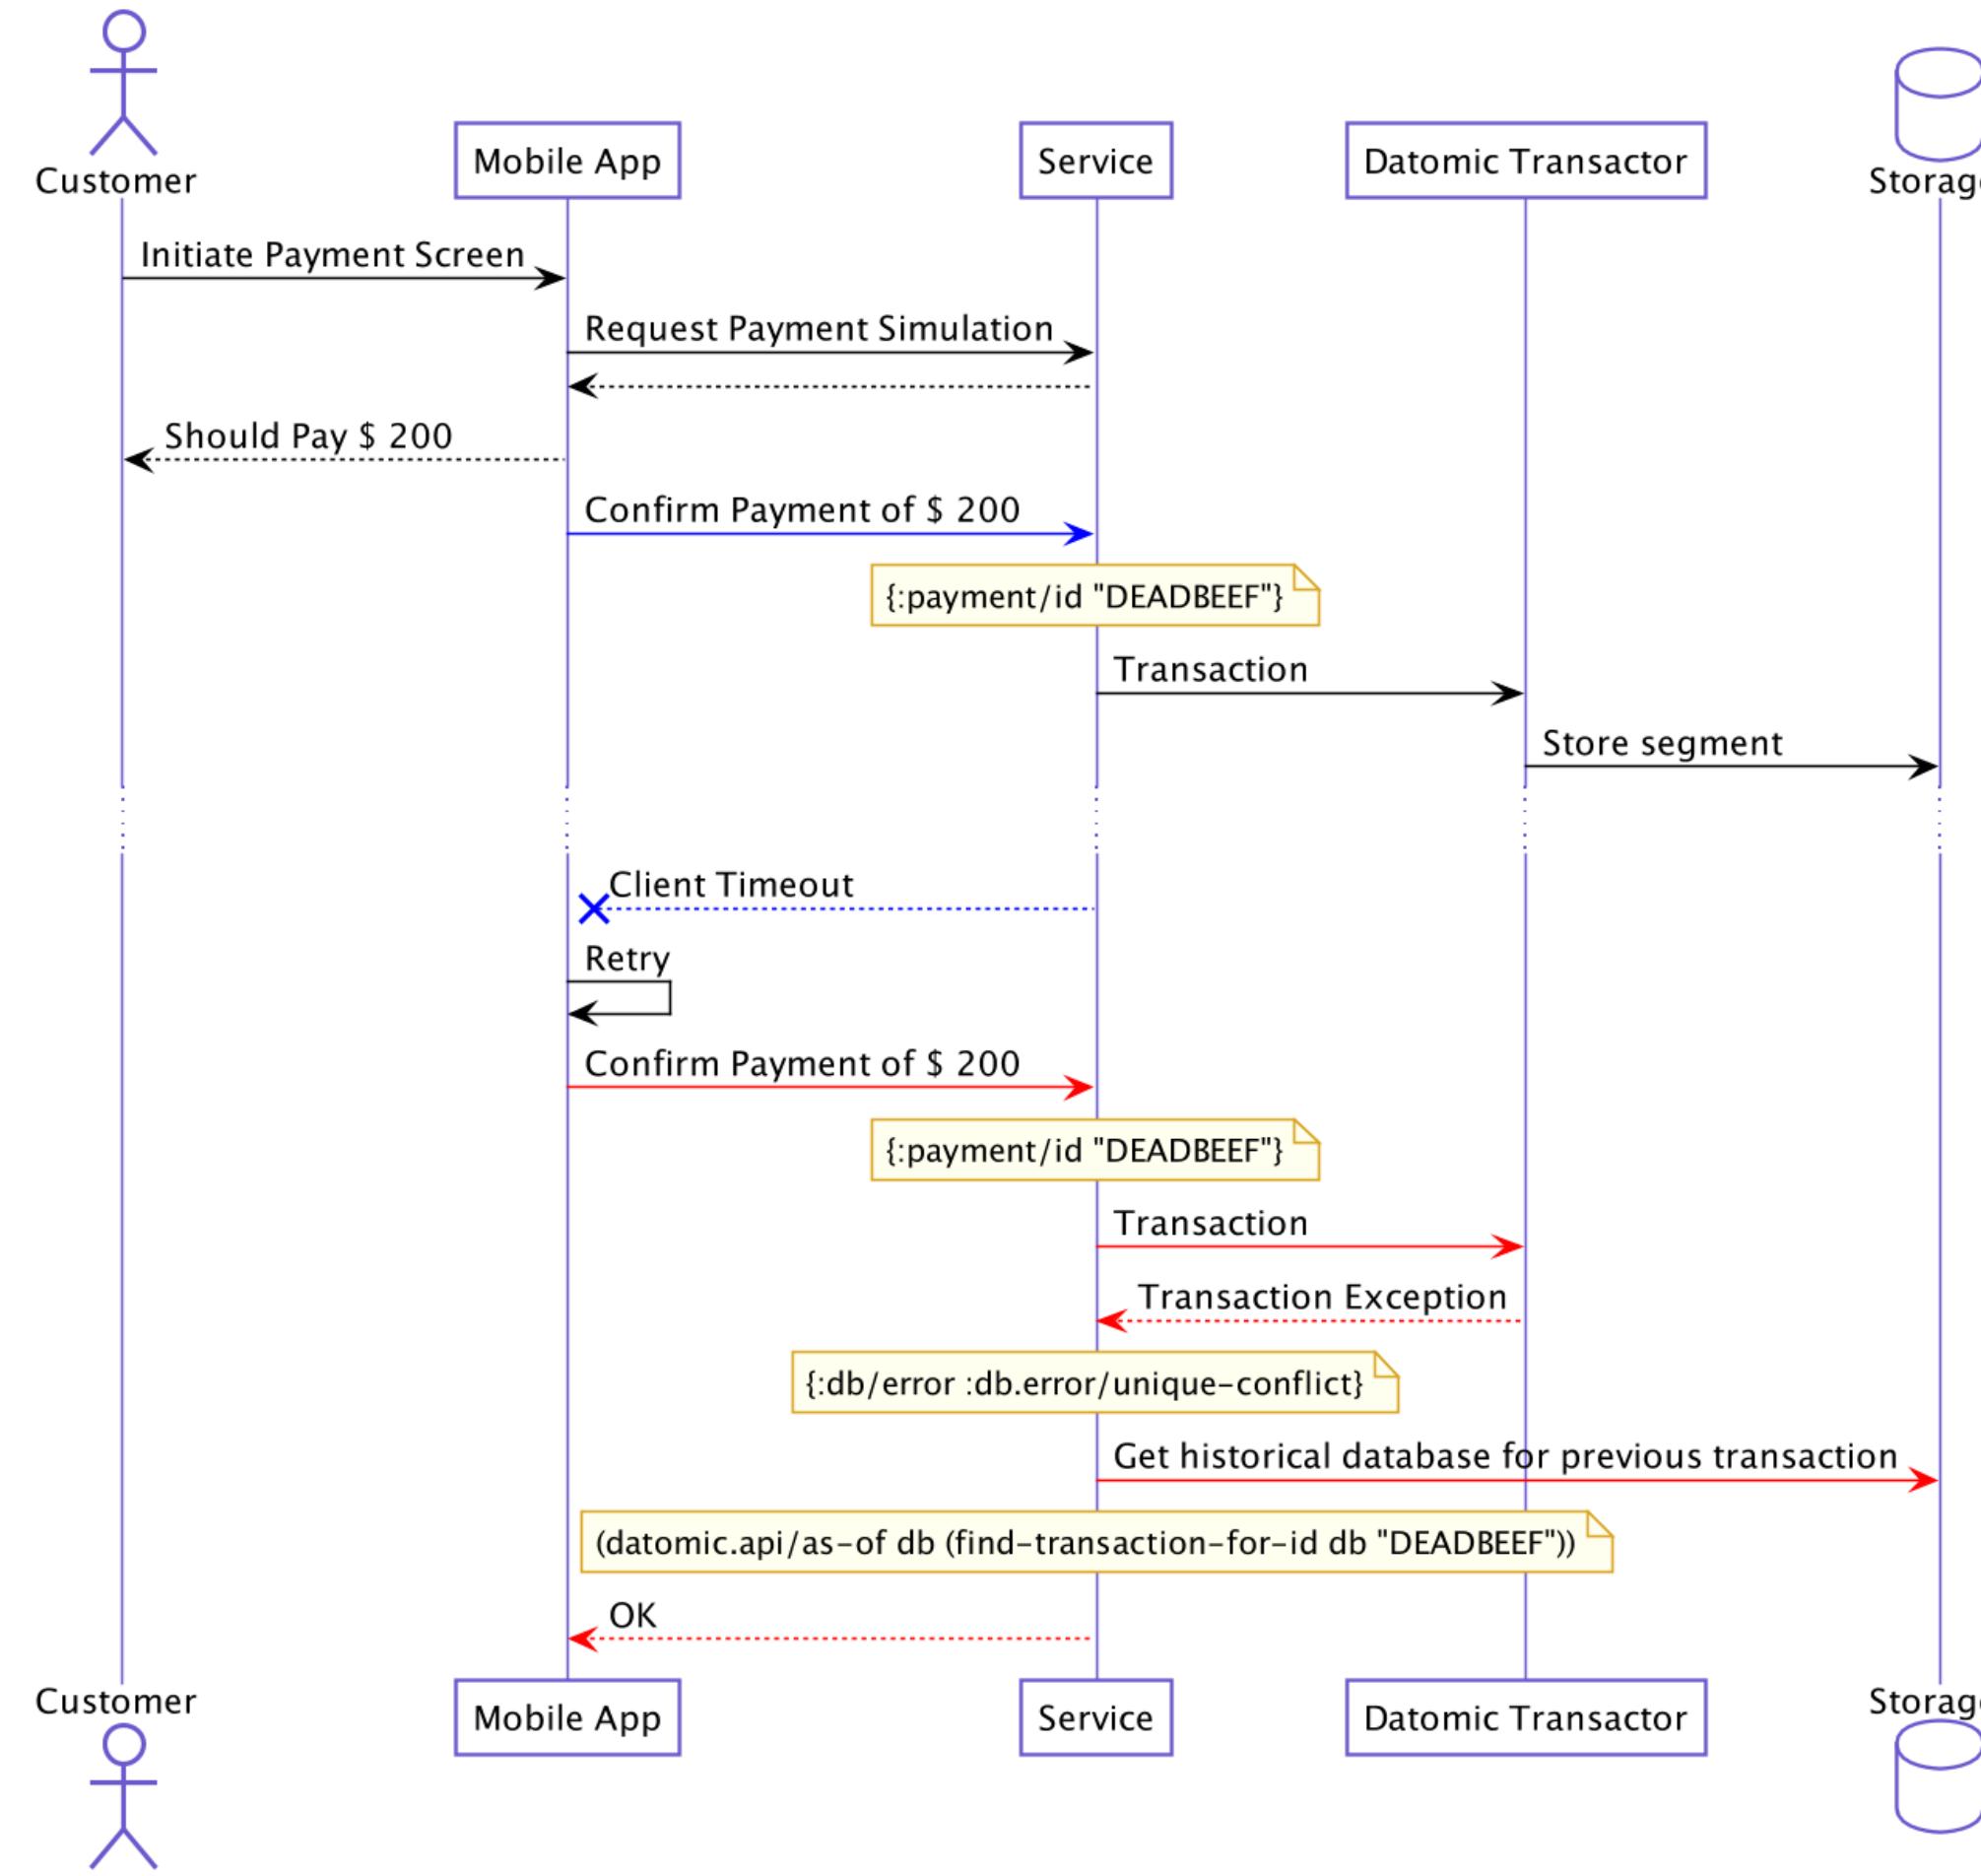
function

Usage: (as-of db t)

Returns the value of the database as of some point t, inclusive.

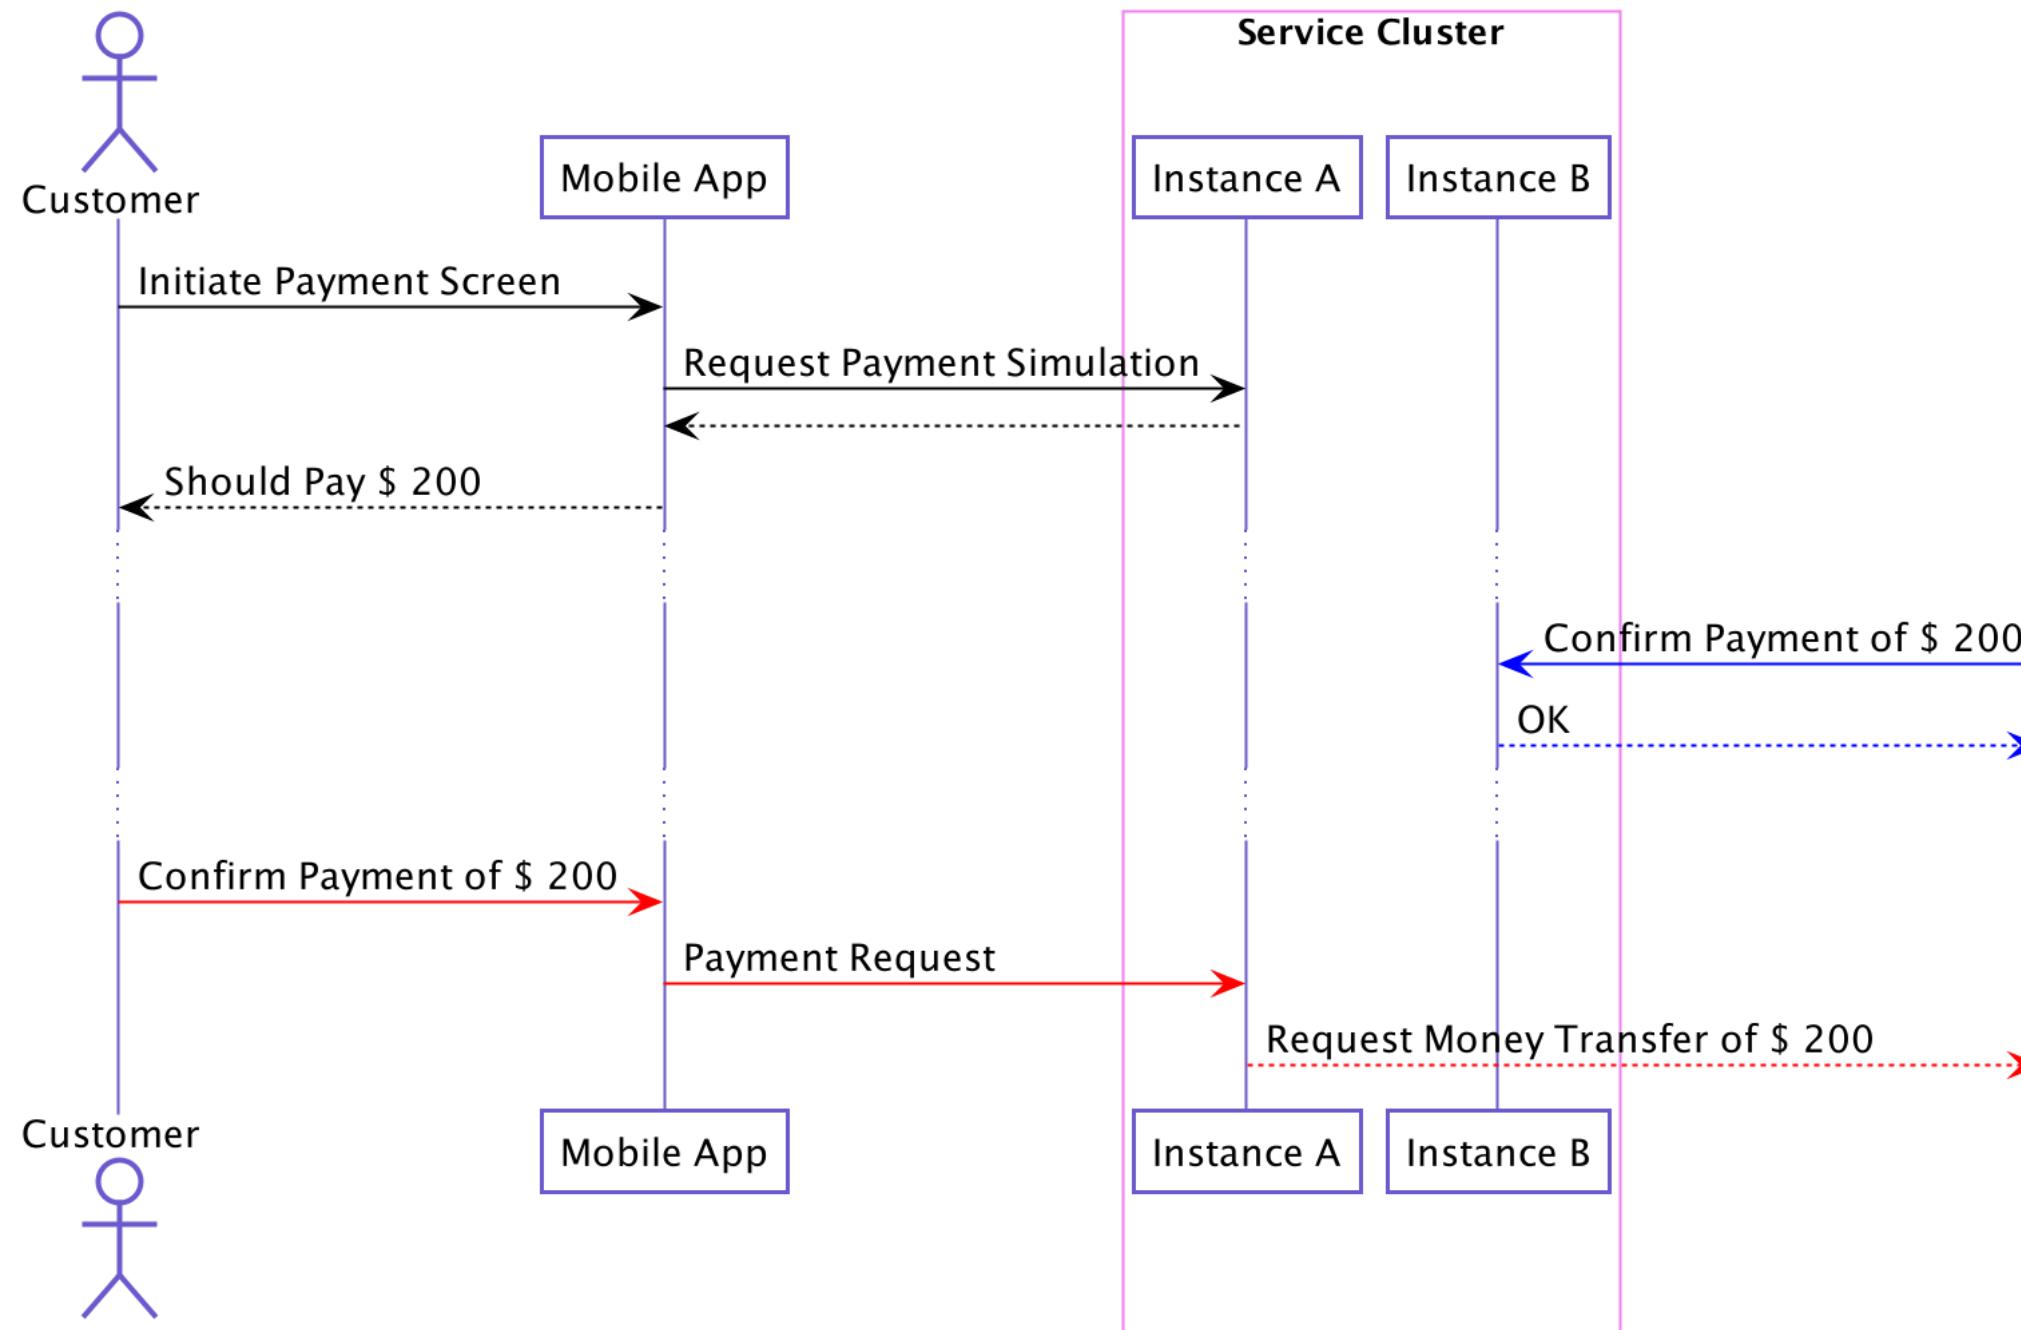
t can be a transaction number, transaction ID, or Date.¹¹

¹¹ <https://docs.datomic.com/on-prem/clojure/index.html#datomic.api/as-of>

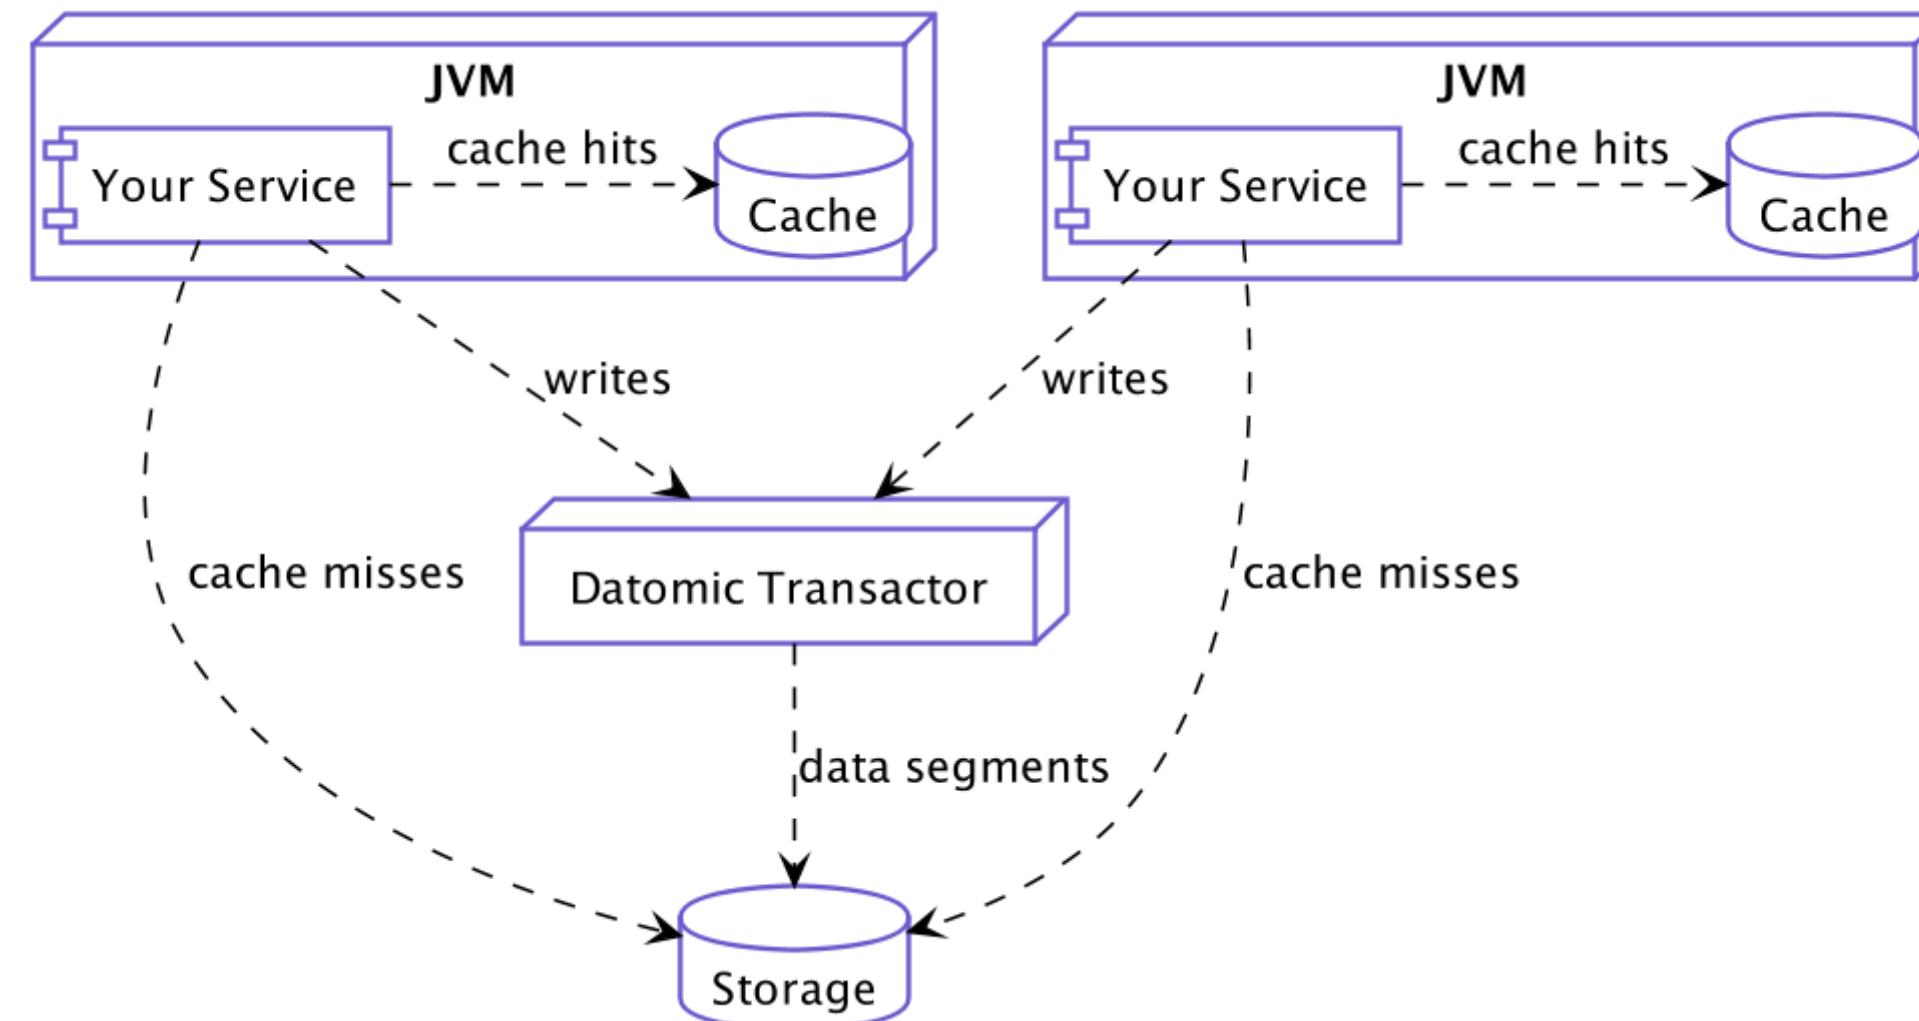


Problem:

Concurrent updates



Quick detour: Datomic's architecture

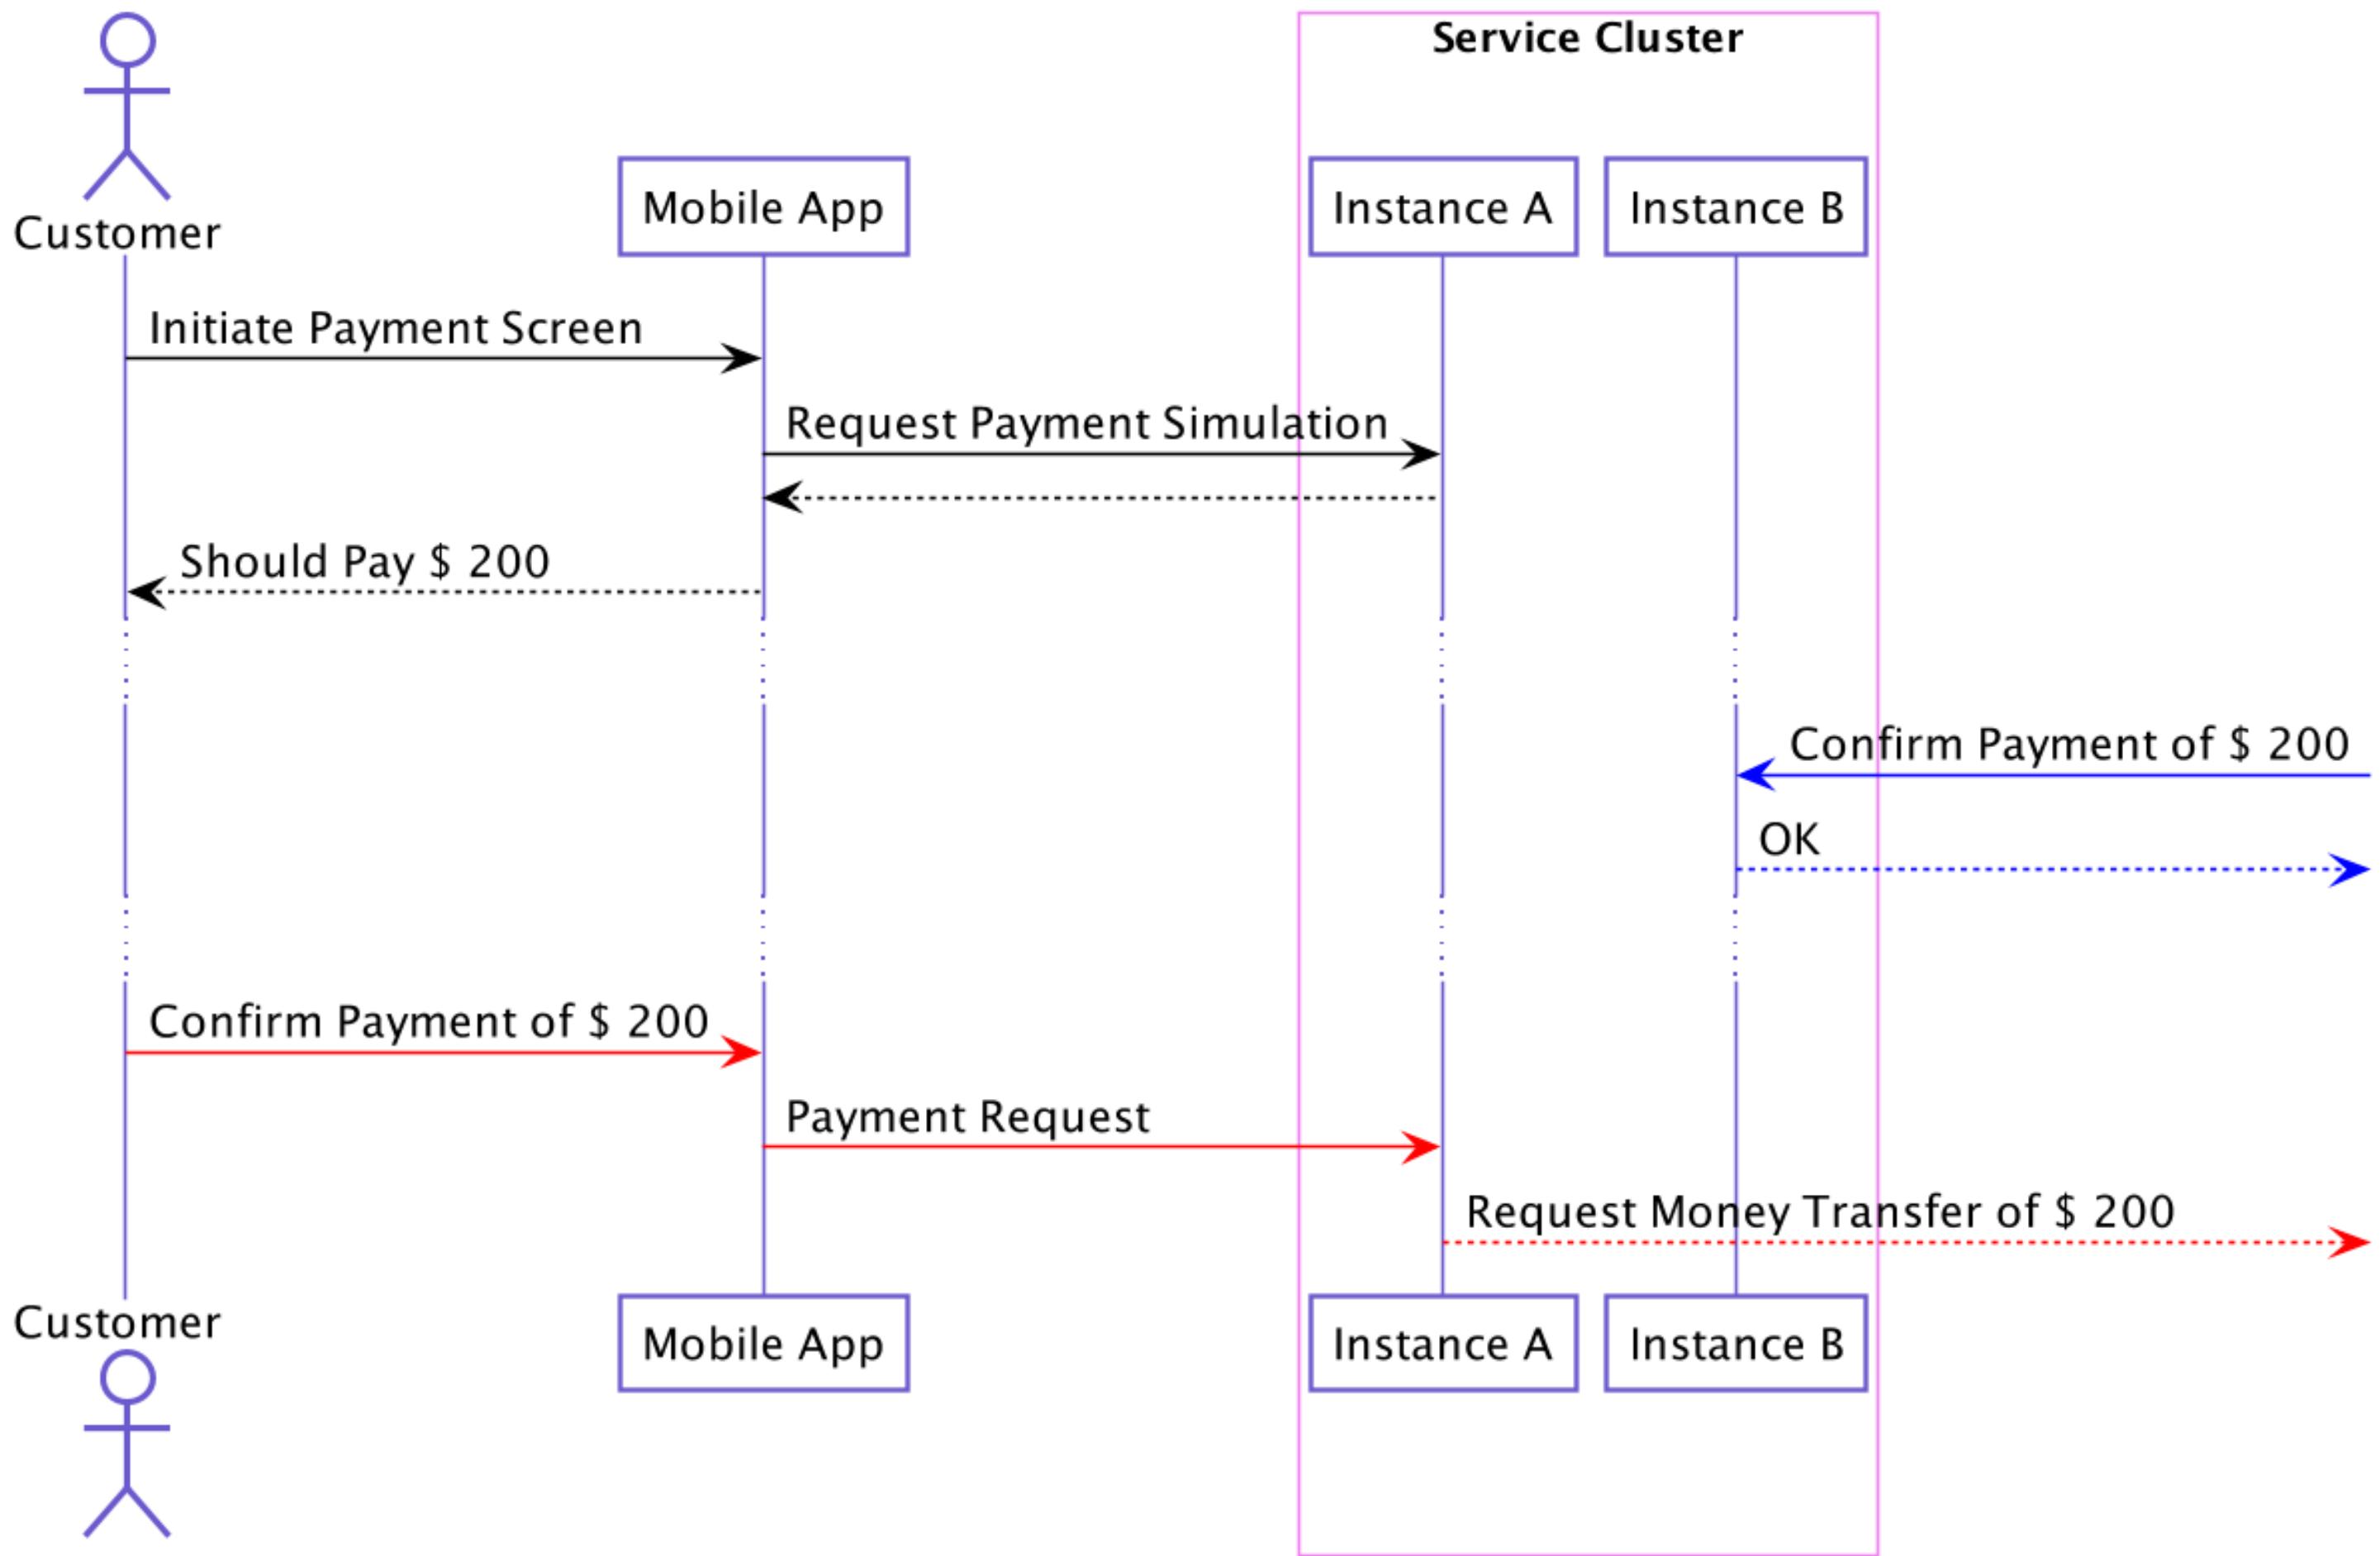


Quick detour: Datomic's architecture

- High-availability for reads
 - Cache locality
 - Cache eviction handled for you automagically
 - Optional: Memcached layer
- ACID transactions

Quick detour: Datomic's architecture

- High-availability for reads
 - Cache locality
 - Trade-off: eventually consistent
 - Cache eviction handled for you automagically
 - Optional: Memcached layer
 - ACID transactions



Idea: Datomic Transaction Functions & HTTP Conditional Headers

:db/cas

The [compare-and-swap](#) function takes four arguments: an entity id, an attribute, an old value, and a new value.

If the entity has the old value for attribute, then the new value will be asserted.

Otherwise, the transaction will abort and throw an exception.⁹

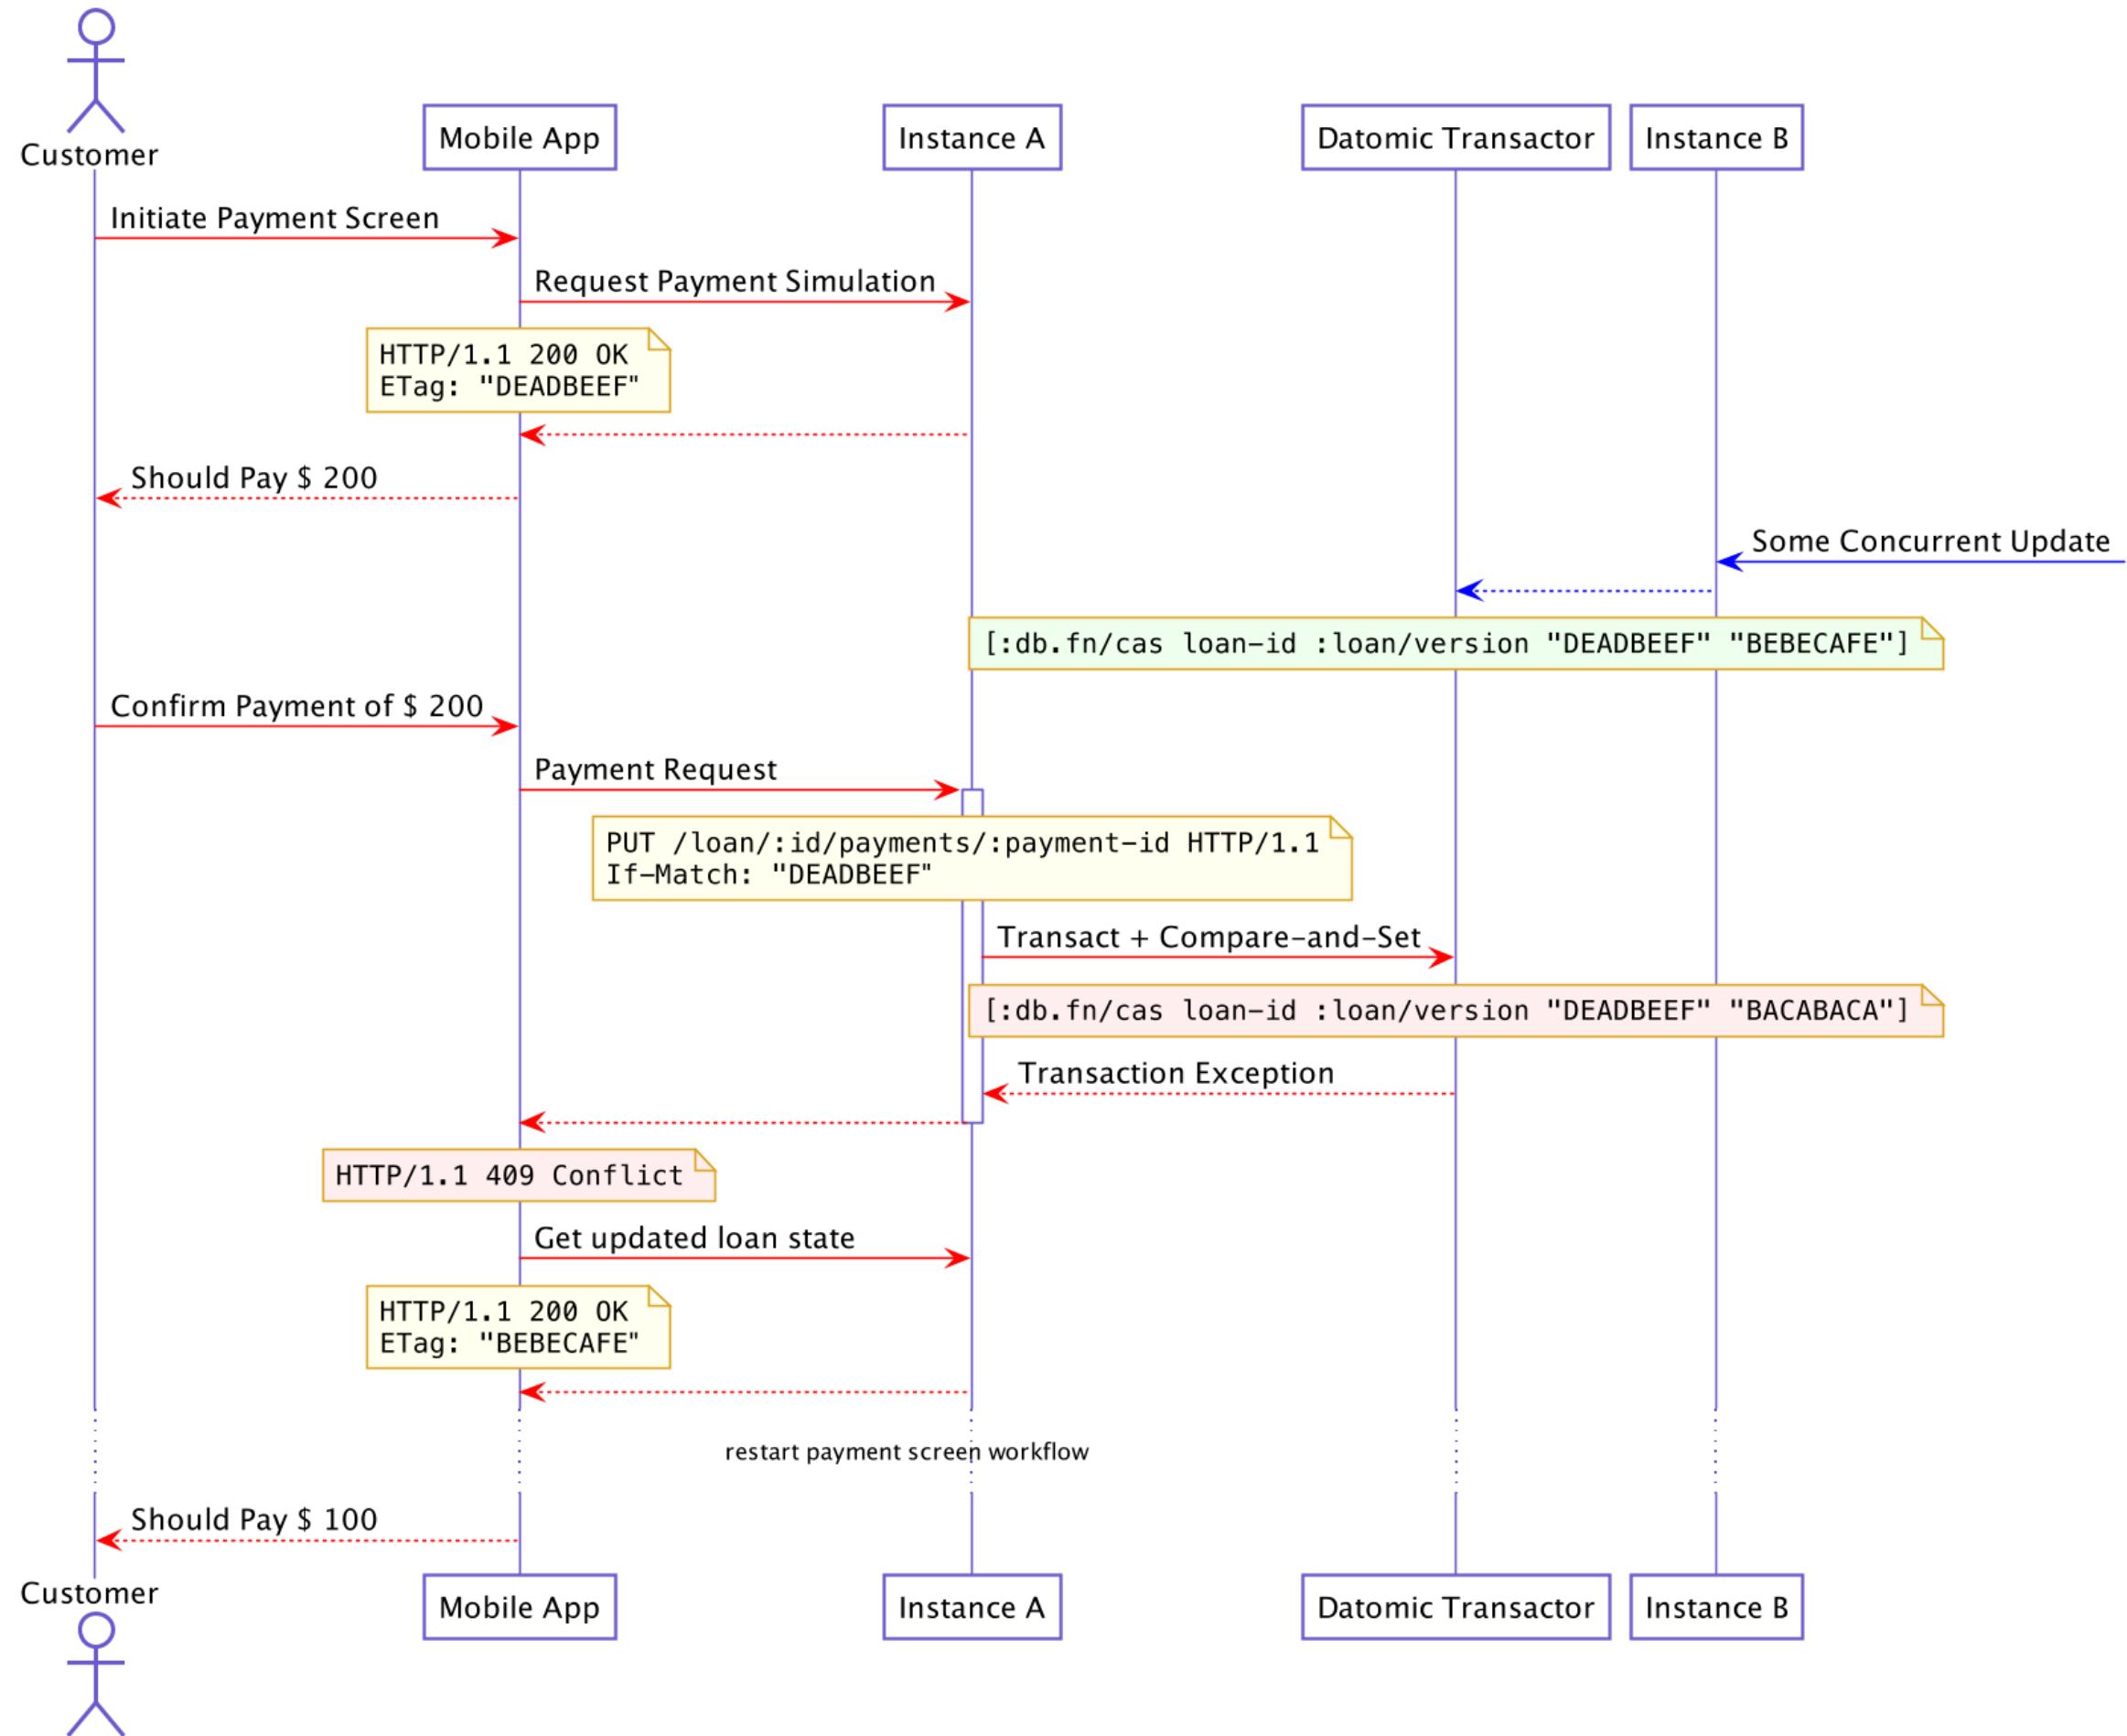
⁹ <https://docs.datomic.com/cloud/transactions/transaction-functions.html#db-cas>

If-Match

The [If-Match](#) HTTP request header makes the request conditional.

For other methods, and in particular for PUT, [If-Match](#) can be used to prevent the lost update problem. It can check if the modification of a resource that the user wants to upload ([will not override another change that has been done since](#) the original resource was fetched. If the request cannot be fulfilled, the 412 (Precondition Failed) response is returned.¹⁰

¹⁰ <https://developer.mozilla.org/en-US/docs/Web/HTTP/Headers/If-Match>



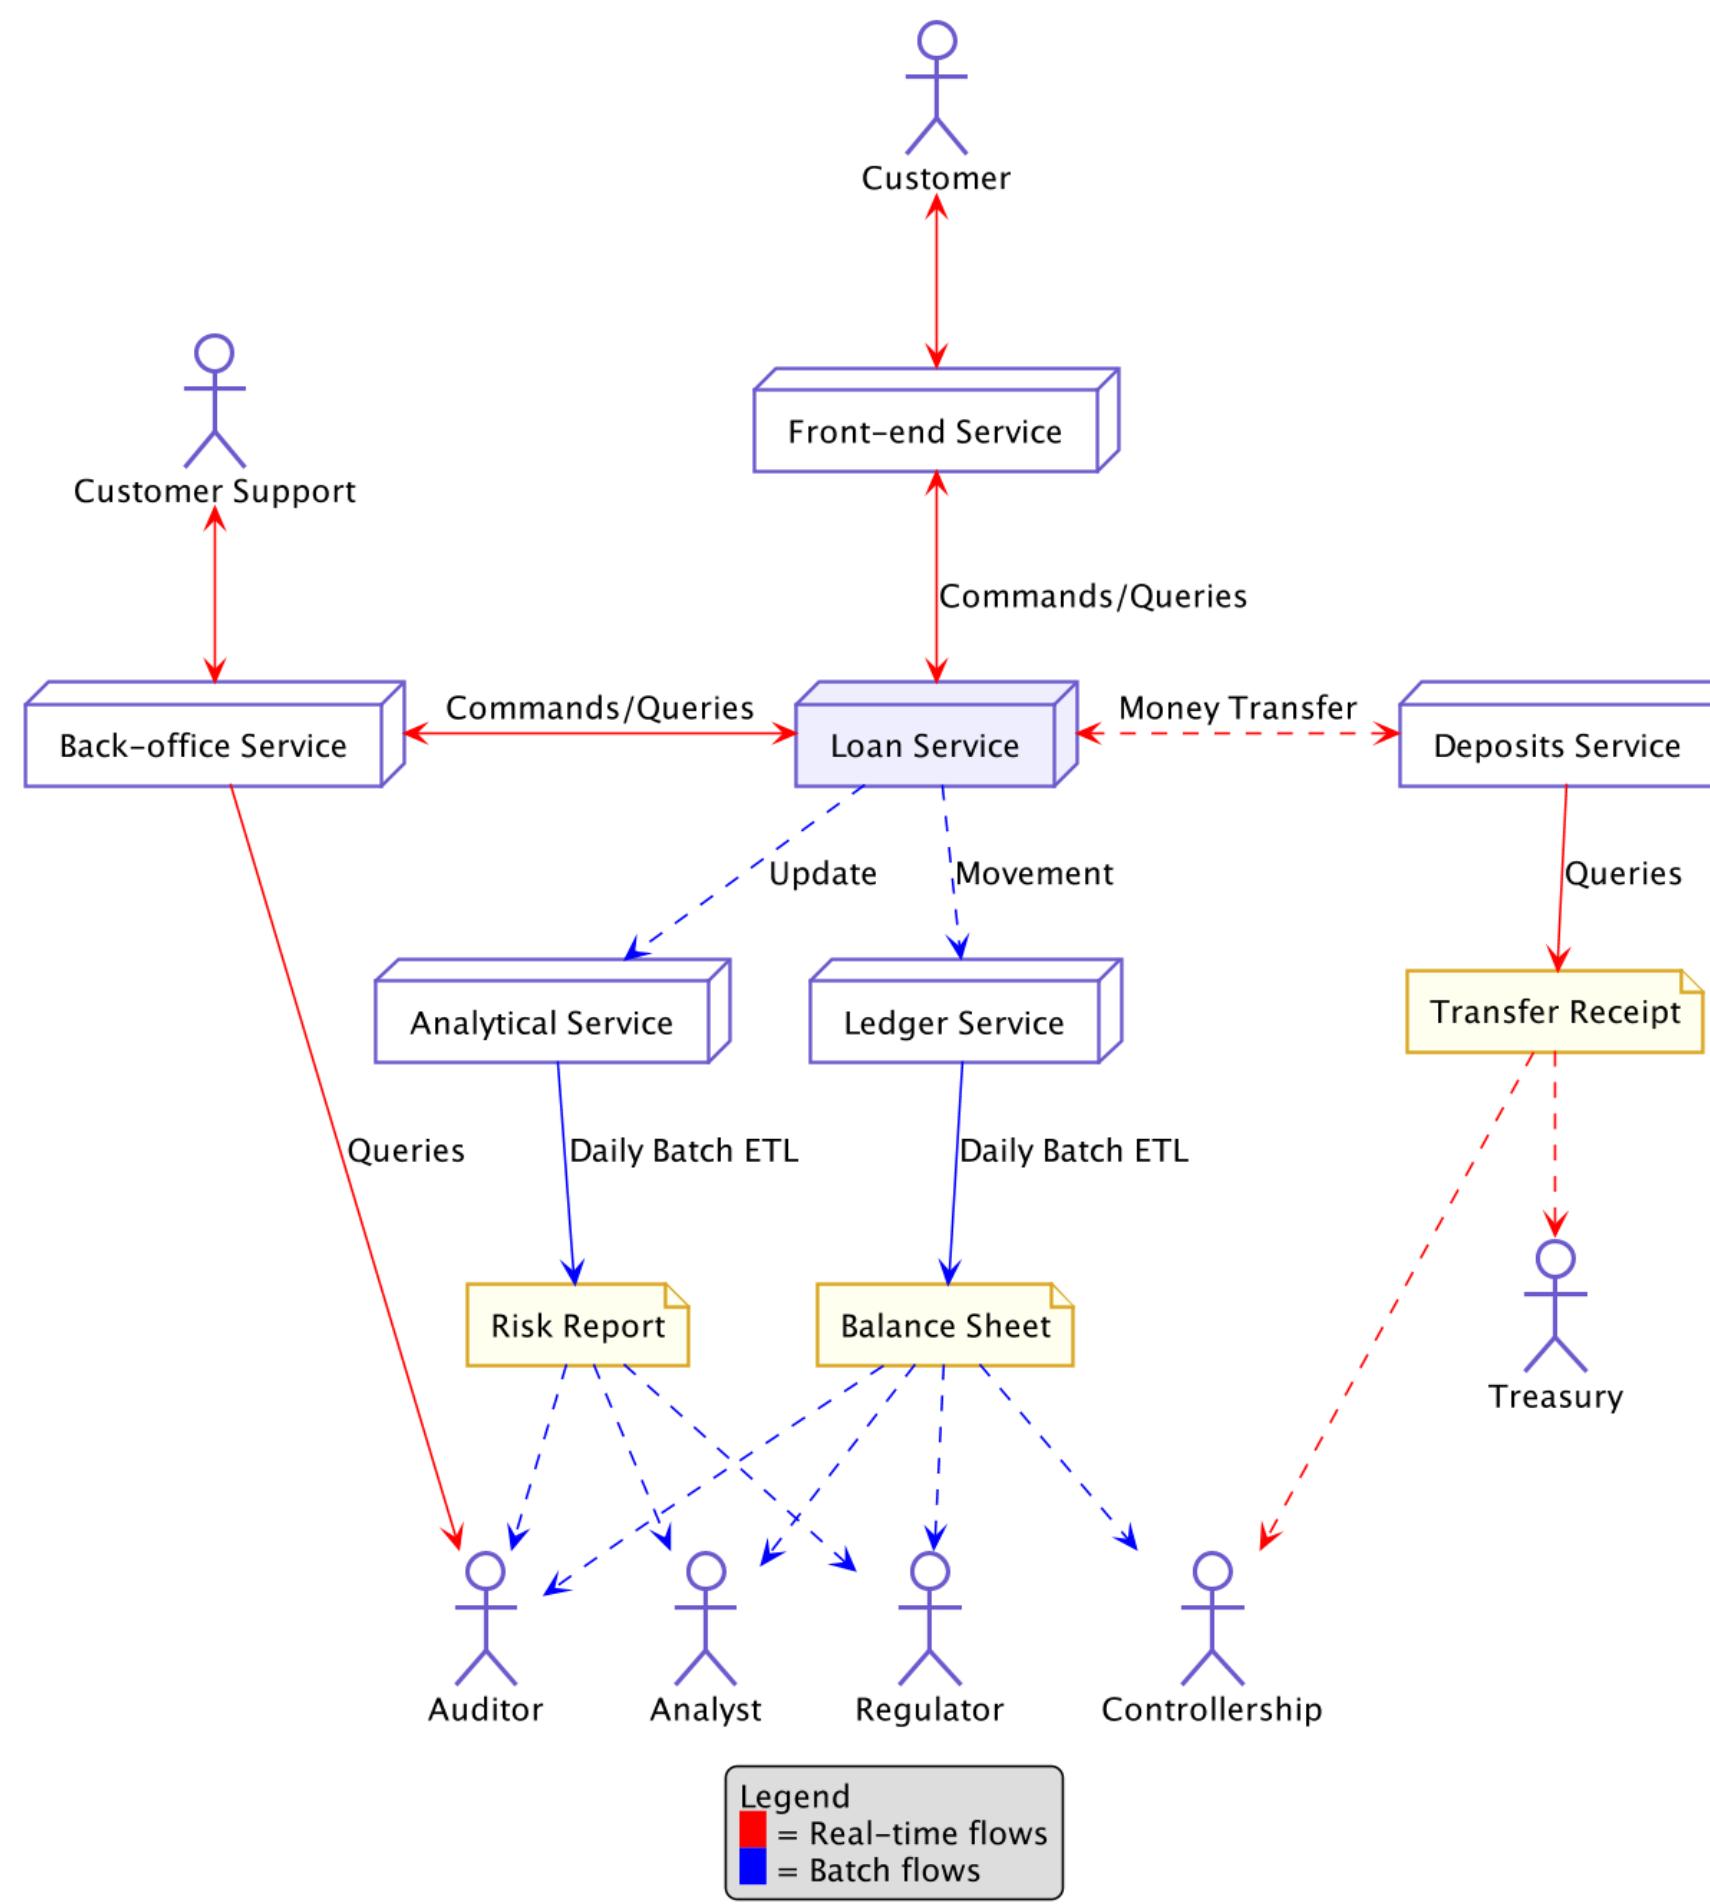
HTTP <-> Datomic Feature Correspondence

Verb	Req Header	Resp Header	Parameter	Datomic API
GET	-	ETag	-	Historical DB (as-of)
PUT	If-Match	-	-	ACID Transactions + Compare-and- Swap (:db.fn/cas)
PUT	-	-	dry-run	Virtual transactions (with)

chapter 3: System

Introducing: the actors

- Customer: Borrows and pays loans.
- Treasury: Manage cash stock for daily operation.
- Controllership: Reconciliation, incorporate product financials on balance sheet, collect taxes.
- Analysts: Understanding customer, product and portfolio behaviors.
- Regulator: Ensure country-level financial system health.
- Auditor: Cross-check data from different sources.



Idea: Reporting Databases

The **operational needs** and the **reporting needs** are often quite **different** - with different requirements from a schema and different data access patterns.

(...)

It's often a wise idea to separate the reporting needs into a reporting database, which takes a copy of the essential operational data but **represents it in a different schema**.⁷

⁷ <https://martinfowler.com/bliki/ReportingDatabase.html>

Data shapes

desired read shape

k/v

row

column

document

graph

Datomic supports via

AVET

EAVT

AEVT

EAVT + pull

VAET

Keeping sanity

- Separate transactional workflows from reporting workflows
 - Denormalized views with the right granularity required by the actors
- Focused core domain services
 - e.g. thin REST API, easy to test
- Auditing is orthogonal
 - Datomic's reified transactions help with that too⁸

⁸ <https://www.youtube.com/watch?v=7lm3K8zVOdY>

Conclusion

Feature

Functional Programming +
Values

Persistent Data Structures

Speculative transactions

Reified Transactions

Data shapes

Ideas

Use Business Language,
Domain Driven Design

Event Sourcing

Command / Query
Responsibility Segregation

Orthogonal Audit

Reporting Databases

Conclusion

- Know your stakeholders and its timing requirements (real-time vs. batch, sync vs. async, etc)
- Fight entropy
- Use separate databases and shapes for reporting
- Keep tooling simple

Thank you!

Henrique Alves hcarvalhoalves@gmail.com

Gustavo Barrancos gbarraconos@gmail.com